

Sehbau: The Software Suite

C. Rasche

<https://github.com/Sehbau>

July 22, 2025

This document describes the Sehbau software suite, a computer vision system that operates with contour and region information. The system distinguishes itself from previous approaches by a much faster and richer feature extraction process. It uses a divisive segmentation procedure, that returns all contours and region boundaries irrespective of contrast and size. The features are thoroughly parameterized and the resulting description can be used for identification and categorization of any image size and any structure; and for determining motion flow. Those recognition processes can be carried out based on the *same* feature extraction output, and not on multiple feature-extraction phases as in other methodologies. This enables to build fast and flexible recognition pipelines. The description can certainly be combined with other methodologies such as Deep Networks and Local Features. The software suite is the ideal substrate for building an active vision system. The most comprehensive description of the system is available under:

<https://www.researchgate.net/publication/391240551>

The software suite comprises a set of program binaries, that carry out different recognition phases such as descriptor extraction, descriptor matching, unsupervised segmentation, retrieval and learning. Examples of how to read the program output are given with Matlab, and partially with Python (version 3.11.9). The code notation is explained in Appendix B.

The introduction explains the architecture for feature extraction and the presently available descriptors (Chapter 1). And it surveys the individual programs, both implemented and planned. The starting point for recognition is the program for *descriptor extraction*, called **dscx**, whose file output is explained in Chapter 2. The remaining chapters explain how to match and deploy the descriptor output and how to further analyze a scene. The chapter on applications proposes how to deploy the suite for specific tasks (Chapter 9).

The software can be found on

<https://github.com/Sehbau/Haupt>

for the following compilations, all 64 bit (x86):

Windows	SEHBAU_win.zip
Ubuntu, 22.04.4 LTS	SEHBAU_ubu.tar.gz, compiled under WSL2
Debian	SEHBAU_deb.tar.gz, compiled under WSL2
Fedora	SEHBAU_fed.tar.gz, compiled under WSL2
...more in the making...	

This document is also available on:

<https://www.researchgate.net/publication/391238505>

and that version will contain the latest updates.

Contents

1	Introduction	1
1.1	Folder Content	1
1.2	Survey of Programs and Usage	2
1.2.1	Principal Programs and Pipeline	2
1.2.2	Introduction to Usage	4
1.3	Structural Description	5
1.3.1	Architecture	5
1.3.2	Feature Extraction	6
1.3.3	Descriptors	7
1.3.4	Texture	10
1.3.5	Saliency	11
1.3.6	Proposals	12
1.3.7	Representation Formats	12
1.4	Program Binaries	14
1.4.1	Computation	14
1.4.2	Learning	14
1.4.3	Utility	15
1.5	Demonstration, Varia	15
1.6	Recognition Pipeline	17
2	Descriptor Extraction [/DescExtr]	18
2.1	Program Use [dscx]	18
2.2	Output Files	20
2.2.1	Description Image (.dsc)	20
2.2.2	Histograms (.hst , .kol)	24
2.2.3	Texture Maps (.txm)	24
2.2.4	Saliency (.slc)	25
2.2.5	Proposals (qbbx , qdsc)	27
2.3	Options and Parameters	29
2.3.1	Architecture	29
2.3.2	Contours	30
2.3.3	Regions	30
2.3.4	Radial Shape	30
2.3.5	Partitioning (Arcs/Straighters)	31
2.3.6	Shape (Arcs & Strs)	32
2.3.7	Texture	32
2.3.8	Utility	32
2.4	Collecting Histograms [h2arr , collhimg]	33

2.5	Generating Vector Files	33
2.5.1	One Image Description [<code>d2vmx</code>] (<code>.vecCnt</code> , <code>.vecRsg</code> , ...)	34
2.5.2	List of Image Descriptions [<code>collvec</code>]	34
3	Matching Vectors [<code>/MtchVec</code>]	35
3.1	Program Use [<code>mvec1</code> , <code>mvecL</code>]	36
3.1.1	Options	36
3.2	Output	37
3.2.1	Program <code>mvec1</code>	37
3.2.2	Program <code>mvecL</code>	38
3.3	Motion Vectors [<code>motvec</code>]	39
4	Matching Histograms [<code>/MtchHst</code>]	40
4.1	Histogram-of-Attributes [<code>mhstL</code>]	40
4.1.1	Program Use	40
4.1.2	Output	41
4.2	Kolumns [<code>mkolL</code>]	41
5	Focus Selection [<code>/FocSel</code>]	42
5.1	Program Use [<code>focdsc1</code>]	43
5.2	Output	43
6	Shape Extraction [<code>/ShpExtr</code>]	45
6.1	Program Use [<code>shpx</code>]	45
6.1.1	Parameters and Options	46
6.2	Output	46
7	Shape Matching [<code>/ShpMtch</code>]	47
7.1	Program Use [<code>mshp</code>]	47
7.2	Output	47
8	Demo Segregation RGB [<code>/DemoSgrRGB</code>]	49
8.1	Program Use [<code>sgrRGB</code>]	49
8.1.1	Options	50
8.2	Output	50
9	Applications	51
9.1	Classification and Identification	51
9.1.1	Clustering	52
9.2	Recognition for Navigation	52
9.3	Other	53
9.3.1	Small Object Recognition	53
9.3.2	Anomaly and Change Detection	54
9.3.3	Collecting Annotations	54
9.3.4	Active Vision	55
9.4	Methodological Fusion	55
9.4.1	Local Features	55
9.4.2	Deep Networks/Learning	56

<i>CONTENTS</i>	iv
A Feature Files	57
A.1 Bounding Boxes (.Bbox)	57
A.2 Contour Endpoints (.CntEpt)	58
A.3 Boundary Information (.BonBbox , .BonAsp , .BonPix)	59
B Notation	60

Chapter 1

Introduction

We firstly overview the content of the software suite (Section 1.1), followed by introducing the principal sequence of programs ([dscx](#), [mvec](#), [focsel](#), [shpx](#), etc.), with which one can perform image, focus (region) and shape matching (Section 1.2). Then we explain some basic parameters of the architecture and the descriptors developed so far (Section 1.3). We mention the complete list of programs in three sections: the main processes, namely descriptor extraction and their matching (Section 1.4.1); the learning programs (Section 1.4.2); and the utility programs that facilitate matching and segmentation (Section 1.4.3). Section 1.5 explains the available demonstration programs; Section 1.6 outlines how the binaries can be deployed in recognition pipelines.

1.1 Folder Content

The folder `/SEHBAU` contains the following directories, with blue denoting program binaries, that exist in those directories:

<code>/DemoBaum</code>	demo for global-to-local segmentation, baumgrau , baumfarb
<code>/DemoPlcRec</code>	demo for place recognition
<code>/DemoSgrRGB</code>	demo for foreground-background segregation, sgrRGB
<code>/Demos</code>	various scripts demonstrating parameter changes
<code>/DescExtr</code>	descriptor extraction, dscx , and conversions to vectors: h2arr , collhimg , d2vmx , collvec
<code>/FocSel</code>	focus of attention: selects descriptors from a region, focsel
<code>/MtchHst</code>	matching attribute histograms, mhst , mkoll
<code>/MtchVec</code>	matching descriptor vectors, mvec , motvec
<code>/ShpExtr</code>	shape extraction for a patch (image), shpx
<code>/ShpMtch</code>	shape matching, mshp
<code>/UtilMb</code>	utility routines for Matlab
<code>/UtilPy</code>	utility routines for Python
<code>globalsSB.m/.py</code>	global variables (Matlab/Python)
<code>exsbAll.m/.py</code>	summary script running all example scripts

The program binaries do not require Matlab or Python to be executed, but we provide utility and example scripts for both languages, to demonstrate how to run the

binaries and how to read their output (standard output or files). The code for both languages lies in various directories with subdirectories named `/UtilMb` and `/UtilPy`, respectively. The Matlab scripts are the most elaborate and we therefore use extension `.m` throughout the documentation to indicate that we talk of a utility routine or a demo script. The python scripts are few in number so far; more are being developed.

Example scripts contain generally the prefix `exsb` (example Sehbau) and exist in various locations; they must be run from their respective folder as we often use relative paths to the binaries and example images. The simplest two example scripts are:

```
DescExtr/exsbDscxSimp.m (.py)    descriptor extraction dscx
MtchVec/exsbMatch.m (.py)        matching with mvec
```

They use the fewest utility routines for reason of clarity. Other scripts utilize convenience functions.

The script in the main folder called `exsbAll.m` runs all example scripts, starting with those two simple scripts. The example scripts for place recognition contain the prefix `plc`. The script `globalsSB.m` provides paths and sets some global variables.

1.2 Survey of Programs and Usage

The principal programs are surveyed first (Section 1.2.1), followed by a short introduction of their usage (Section 1.2.2).

1.2.1 Principal Programs and Pipeline

To launch recognition, we firstly carry out feature extraction and description for the entire image, executed with program `dscx`. Then we use that image description to match it with stored descriptions, carried out with program `mvec`.

```
dscx    →    mvec
```

Then we start focusing on certain parts of the image by using processes such as focus selection and shape description, executed with programs `focsel` and `shpx`, resp. Those two processes can be regarded as attentional shifts. The output of focus selection can be matched again with program `mvec`:

```
focsel    →    mvec
```

The output of shape description is matched with program `mshp`:

```
shpx      →    mshp
```

We elaborate on the individual processes:

- `dscx` [descriptor extraction]: the binary outputs the features and descriptors into a number of different files of which the three main ones are:
 - the *description* file with extension `.dsc`, containing the descriptor attributes with which one can span a multi-dimensional space, useful for identification.

- the *histogram* file with extension `.hst`, expressing the attributes as histograms, useful for fast classification with a 'traditional' classifier (Linear Discriminant Analysis [LDA], SVM, RandomForest [RF], etc.).
- the *saliency* file with extension `.slc`, containing scene statistics, some object proposals and texture information (Section 1.3.5). This information can be used to decide where to apply focus selection and shape description. And it can be used for visual orienting, for example for auto-focusing, zooming, saccading (change of camera direction), etc.
- `mvec` [matching vectors]: matches the descriptors as outputted by the program `dscx` or `focsel`, and returns metric measurements for various types of descriptors. This is useful for identification of structure. The use of this binary will be explained in Chapter 3.
- `focsel` [focus selection]: extracts the description of a desired region, a so-called *focus*, from the description file as generated by `dscx`. The region is defined by the user as a bounding box and can outline an object proposal or part proposal, ie. obtained from the saliency file; or it can be an annotation. The bounding box can be of arbitrary dimension and size. `focsel` extracts both, vectors and histograms, extensions `.dsf` and `.hsf`, resp. The vectors can be matched with other focii with the program `mvec`. To be further detailed in Chapter 5.
- `shpx/mshp` [shape extraction and matching]: refines the segmentation of a shape and saves it to file with extension `.shp`, which then can be used for matching (Chapters 6 and 7). This shape can be any silhouette, be it the letter of some text in the wild, or an object with homogenous color, or scene part.

With that program survey, we can now refine the above pipeline. Since the deployment of vector matching - using `mvec` - is a relatively costly process, it makes sense to subselect candidates by firstly classifying the histogram output, process `Clsf`, and then to apply `mvec` on the identified subset of representations:

```
dscx → Clsf(.hst) → mvec(.dsc)
```

This multi-stage (cascaded) process can be carried out with a trained classifier (LDA, SVM, RF, etc.), which makes sense if we assume a clearly defined category. Or it can be based on histogram matching only, if the goal is to identify a structure. For the former there exists enough software; for the latter we provide a separate program called `mhst` (Chapter 4).

Focus selection can be based on the output provided by the saliency file, or any output of `dscx`. For matching, we can again make a selection based on classification:

```
focsel(.slc | .dsc) → Clsf(.hsf) → mvec(.dsf)
```

Focus selection allows applying tailored representations, without performing a complete feature extraction and description. It is useful in particular, if we wish to analyze structure containing contour information or texture. If the goal is to analyze rather a shape silhouette in more detail, then we apply shape description and matching:

```
shpx(.slc) → mshp
```


Or one can apply both processes to the same patch. We will refine the pipelines even more, after we have introduced the full set of programs (Section 1.6). Next we illustrate the usage of those programs.

1.2.2 Introduction to Usage

The following examples give an idea of how to provide the arguments to the program binaries.

The task is to compare two images named `imgA.jpg` and `imgB.jpg`. Firstly, we generate the descriptors and provide a filename as output, in this case single letters `A` and `B` (both in directory `/Desc`):

```
> dscx imgA.jpg Desc/A
> dscx imgB.jpg Desc/B
```

This will write the description files called `A.dsc` and `B.dsc` to directory `/Desc`. In a first round we compare the images at once, for which we feed the description files as arguments to program `mvec`,

```
> mvec Desc/A.dsc Desc/B.dsc
```

which returns dissimilarity and similarity metrics, either as standard out (`stdout`) or as file.

In a second round, we compare two different regions, for which we now deploy `focsel`. We select the upper left quadrant as bounding box, `0 128 0 128` (assuming image sizes are both `256x256`):

```
> focsel dscA.dsc 0 128 0 128 focAupplef
```

This will write the subset of descriptors to file `focAupplef.dsf`. We extract an equally sized region from image B and call it `focBsomewhere` (operation not formulated here). Then we match those two regions:


```
> mvec focAupplef.dsf focBsomewhere.dsf
```

which again returns the metric measurements. We integrate the results ad libitum.

The programs allow to take lists as input, ie. one can specify a text file as input that contains file names or bounding boxes. The details of that will follow in the chapters.

1.3 Structural Description

The structural description is based on contours and regions. The extraction of those is called *feature extraction*, resulting in lists of contours and regions (Section 1.3.2). Those features are then partitioned, parameterized and integrated, which is referred to as *feature description*. The output of feature description consists of so-called *descriptors*, explained in Section 1.3.3.



The outputted description is quite rich and can be deployed in various ways, ie. selected and interpreted according to the specific task (Section 1.3.7). And it can be used to describe scene textures (Section 1.3.4) and naturally is suitable for saliency and proposals (Sections 1.3.5 and 1.3.6, respectively). This entire extraction and description process is termed *descriptor extraction*, hence the program name **dscx**. Program **dscx** outputs the descriptors and texture blobs, but also bare region information. We now firstly explain what architectures are available to extract the features from.

1.3.1 Architecture

Contour and region features are extracted from an image space. Two spaces are available: a pyramid and a (cubic) scale space. A space consists of a stack of **nLev** levels. We firstly explain the architecture and parameters for the pyramid, then the ones for the scale space.

Pyramid

The following schematic shows a pyramid made of four levels, **nLev=4**, using zero-indexing:

```
lev 3      --
lev 2      ----
lev 1      -----
lev 0      -----
```

The bottom level of the pyramid, **lev 0**, holds the original image resolution. Higher pyramid levels are generated by downsampling with an integer factor equal two. Downsampling continues until the map is equal 16 pixels; or just larger, for the smaller side length. In the depicted schema, level equal 3 would be the top level and would correspond to the pyramid of a 128x128 pixel image. The number of levels can be set as command argument using double dash to specify a long option (**--nLev**):

```
> dscx imgA.jpg /dscA --nLev 2
```

The parameter can also be set by file, the details of that follow in later sections. For each level of the pyramid, contours and regions are extracted, which results in a vast set of features. For efficient recognition, some sort of selection must take place. We firstly introduce feature extraction and its selection parameters (Section 1.3.2), followed by the descriptors developed and their selection procedures (Section 1.3.3).

There are no particular image preprocessing algorithms carried out in this program, as I have never observed any consistent improvement of using for instance smoothing or low-pass filtering. In fact, I have the impression it is detrimental sometimes. But there may be tasks or databases where image filtering could be useful. In that case, it is best to carry out the preprocessing separately (beforehand) and then feed the preprocessed image to `dscx`.

Scale Space

The pyramid architecture is suitable for the fast analysis of arbitrary image content. It is the default architecture. If we know what type of input we face, then a scale space may be better suited, in particular when subtle differences need to be discriminated. We can set the architecture to be a scale space by long option `--is`:

```
> dscx imgA.jpg /dscA --is 2
```

and specifying a value of two. A value equal one specifies the pyramid. A scale space takes a little bit longer to compute, as no reduction in space occurs. And it generates more features for the same reason.

An example script comparing the output of the two spaces is given in `exsbImgSpaces.m` in directory `/demos`. The plotting scripts mentioned in later sections, are setup to demonstrate the output for the pyramid, but can be easily modified for the output of a scale space.

There exist more parameters that can be considered part of the architecture, such as the dimensionality for spatial histogramming, but those will be mentioned later.

1.3.2 Feature Extraction

Contours `cnt`, `rre`

Three types of contours are extracted, namely ridge, river and edge contours, sometimes abbreviated as `RRE` or `rre`. They are kept separately initially and we depict that as three separate pyramids. We use symbols `.` and `|` to symbolically express that the pyramid content is individual to an, meaning the levels have different list lengths with different segments lengths (no actual correspondence is depicted with those symbols):

	ridges	rivers	edges
lev 3	.		..
lev 2	.	.	
lev 1	.	.	.
lev 0

The threshold parameter for accepting a pixel value as a contour pixel, is set with `Cnt.minCtr` by file or `--cntMinCtr` by long option. By default the value is 0.05 and is relative to the maximum value of the range image (taken with a 3x3 neighborhood). Directory `/DescExtr/UtilMb/Examples` contains an example script for the detection of these three types of pixels, called `e_CntMap.m`.

Regions `reg`

Regions are detected by a hierarchical thresholding process, whose tree `depth` is typically set to value equal three for images up to ca. 400 x 500 pixels.

The thresholding process is applied to each level. The following illustrates the map output for an architecture with four levels and depth equal three:

	depth 0	depth 1	depth 2
lev 3	--	--	--
lev 2	----	----	----
lev 1	-----	-----	-----
lev 0	-----	-----	-----

The creation of small regions can be controlled by parameter `Reg.minPixNode` (or `--regMinPixNode` by long option), more details in Section 2.3.3. There is no contrast threshold applied here.

The folder `/DemoBaum` contains programs that demonstrate the output for the original resolution, that is `lev=0`, see Section 1.5.

Each of those maps is then analyzed for connected components and their boundaries are extracted:

`regions` \longrightarrow `boundaries`

The boundaries are concatenated across depth (per level), resulting in a pyramid structure as depicted already above for contours. The boundary pixels can be saved to a file with extension `BonPix` (Section A.3).

The divisive region segmentation process is capable of returning regions of lowest contrast possible. This sensitivity results in a vast feature output that is useful for search scenarios, such as defect detection of material surfaces, anomaly detection, detection of occluded objects or of low-contrast objects, etc. The full output can be observed with the files explained in Section A.3.

But for most recognition tasks it requires some sort of selection, based on parameters such as contrast, minimum size, minimum spacing, etc. Some of this selection is better performed on the parameterized features, the descriptors. The default parameters for selection are set such, that one can perform decent classification and identification of the images as appearing in image collections with daily scenes. For a specific task, one might have to adjust some of the parameters, which will be discussed throughout the explanations and be subject when discussing applications (Chapter 9).

1.3.3 Descriptors

The features are parameterized directly, but some are also partitioned and then further parameterized. We first give an overview of those processes, and then introduce the individual descriptors.

Overview

The contour features are directly parameterized without any partitioning, which in turn is used for creating a skeleton subset, a texture description, as well as groups of contours, called bundles:

$$\text{rre feat.} \xrightarrow[\text{description}]{\text{parametric}} \text{rre attributes} \xrightarrow{\text{analysis}} \text{skeleton, texture, bundles}$$

The boundaries themselves will be used for a simple description, called radial-shape descriptor or sometimes just radial descriptor:

$$\text{boundaries} \xrightarrow[\text{description}]{\text{direct}} \text{radial-shape descriptor}$$

In addition, boundaries are also partitioned into curved segments (arcs) and straighter segments (straighters):

$$\text{boundaries} \xrightarrow[\text{partit.}]{\text{curve}} \text{arcs \& straighters} \xrightarrow[\text{analysis}]{\text{statistical}} \text{shape} \xrightarrow[\text{analysis}]{\text{geometric}} \text{tetragon}$$

which in turn are used to create a shape description based on arc and straighter segments. The shape description in turn is used to obtain a geometrically more precise description called tetrasons.

The parametric description is typically carried out for all features, and the resulting lists of descriptors are also called the *full set*. When we subselect features for a specific task, then the resulting subset is called *skeleton* in our publications, but we use for arcs and straighters another term as well, for clarity. We here mention only the key parameters for this selection. The complete list of parametric options will be provided in later sections, as well as the individual attributes. There are currently 7 descriptors, of which the following four are considered basic descriptor types:

- **Contour (cnt)** : describes the ridge, river and edge contours (RRE). An individual contour is described by its length and angular orientation. The use of all three types constitute the full set, also called *RRE set*. Based on this RRE set, we derive three types of descriptions:
 - Skeleton: a selected set of longer contours, called *skeleton* sometimes. The skeleton is saved to the description file and appears as **ACNT** when loading (a list of levels).
 - Bundle: groups of contours, to be further introduced below.
 - Texture: a texture analysis based on the orientation angle of segments, to be introduced in Section 1.3.4.

The RRE set is *not* saved by default, due to their size. It can be saved by turning on long option **--saveRRE**, in which case it is saved to a separate file with extension **dscRRE**.

The reduction from the RRE set to the skeleton set occurs by a global-to-local selection procedure, whose parameters are minimum spacing and minimum contour length, see **sklMinSpc** and **sklMinLen** of Section 2.3.2.

- **Radial shape (rsg)** : describes the region boundary using a radial signature. The description occurs for boundaries of minimum contrast and minimum size, see **rsgMinPix** and **rsgMinCtr** of Section 2.3.4.
- **Arc segment (arc)** : describes a curved boundary segment obtained from partitioning the region boundaries. The full set of arc segments is reduced to a skeleton called here *gerüst* (scaffold). Only this subset is saved to file. The full

set can be saved as well, by turning on long option `--saveCVP`; it is then written to a separate file with extension `dscCVP`.

- **Straighter segment (`str`)** : describes a straighter boundary segment obtained from partitioning the region boundaries. The straighter segments are the segments that lie between the arc segments. Analogous to arcs, only the selected set is saved, the full set can be saved together with the full set of arcs.

These basic descriptors alone provide relatively good categorization and identification performance. But for interpreting our environment more efficiently, we form more complex descriptions:


- **Shape abstraction (`shp`)** : a description based on the segment statistics of arcs and straighters for each individual shape. It contains dozens of attributes. By default, this shape description takes place only for regions that are fully inside the image, that is not touching any image border. There exist scenes that lack any inside shapes, such as landscape scenes, photos of smooth surfaces, etc. In that case, it might be worthwhile including shapes that touch the image border. This can be regulated with parameter `bordTouches`. By default it is set to 0 and ignores any regions, that touch an image border. With value equal 1, a shape can touch one image side; with value equal 2, two sides; with value equal 3 three sides; and with value equal 4 all four sides. The more image sides a shape touches, the more likely it represents a background region; or an object very close to the camera.
- **Tetragon descriptor (`ttg`)** : a description focusing on shapes that contain at least two loosely parallel straighter segments, whose sides appear to form a tetragon that is loosely aligned with either vertical or horizontal image axis. The idea is to describe in particular horizontal and vertical structures that are ubiquitous in scenes and we therefore determine the axis of such a tetragon. The tetragon is a subset of the shape descriptors (`shp`), but with a more refined parameterization.
- **Bundle descriptor (`bnd`)** : a description based on groups of RRE contours. Groups are detected when determining the skeleton set and therefore show the minimum length specified by `sklMinLen`. Groups of shorter segments are better expressed with the texture analysis.

With the attributes of a descriptor type we span a multi-dimensional *attribute space*, that is suitable for identification of structure. Each descriptor type is determined for each level of the image space (pyramid or scale space), which is called the *descriptor image space*, or in short *descriptor space*. The program `dscx` outputs the descriptor spaces to the description file (with extension `dsc`), placed into directory `/Desc` in our code examples. This is summarized in Section 2.2.1. It also outputs them as bins to a `dsb` file for building histograms.

Each descriptor type comes with position and angular attributes. The position is often the (normalized) coordinates of the center pixel of the feature, or it is calculated from some keypoints. The angle attribute describes the orientation of the feature in the image plane. More explanations on the attributes will follow when introducing the individual programs.

1.3.4 Texture

The texture description is based on the full set of contours - the RRE set - of the first level of the image space (`lev=0`). The description occurs in several steps. In a first step we observe the statistics for the segments lying within a window. One type of statistics are histograms of the contour's attributes, called *kolumns*. Then we analyze the statistics in such kolumn histograms, thereof forming a variety of texture maps. Those maps are then described globally by some simple statistics, and also by region, identifying texture *blobs*.



The kolumns and texture maps are saved to individual files; the statistics of the texture maps and their blobs to the saliency file (`.slc`) introduced previously.

Kolumns

Kolumns are formed similar to the technique of spatial histogramming of features, as applied in the local feature approach. For a small rectangular window of the image, ie. 16x16 pixels, a 4-bin histogram of the contours' orientation angles is formed, called a kolumn histogram, or simply kolumn. We also form kolumns with the length attribute. The two types are referred to as ori-kolumns and length-kolumns, respectively.

The windows overlap by half their window side. For example for a 256x256 pixel image we choose a 16x16 pixel window: the overlap is then 8 pixels, and the resulting spatial dimensionality measures 32x32 kolumns. Kolumns are saved to a file with extension `.kol`, if flag `--saveKol` is set. By default the window measures 16x16 pixels; its size can be changed by long option `--txws` (Section 2.3.7).

The deployment of the kolumns will be discussed alongside the use of histograms of descriptors, ie. when introducing the representation format (Section 1.3.7) or the file output (Section 2.2.2).

Texture Maps

Each kolumn is analyzed for its statistics and the derived parameteric values are placed back into separate maps of size 32x32 (256x256 image; window size equal 16 pixels). The simplest two measures are the count of contour segments and the lack of any segments in the window, abbreviated `num` and `blk`. When contours are present (in a window), then five orientation parameters are determined, that we also call texture *biases*:

- `Num`: total number of contours present in the window, also called the numerosity bias.
- `Blk`: a window is considered blank (void), if three or fewer contour segments are present. This allows detecting sky region, water surface (with no reflection), or any region void of texture, even if there is an 'errand' streak present.
- `Nil`: degree of 'no' orientation (nil dominance). Often corresponds to foliage.
- `Vrt`: degree of vertical orientation (dominance). Often corresponds to texture in natural scenes, ie. grass, crops, stems, etc.

- **Hor**: horizontal orientation dominance, for example uneven water surface; objects arranged with increasing (spatial) distance, such as a column of cars parked along a street.
- **Axi**: axial orientation dominance, that is vertical and horizontal (approximately) equally present, ie. windows in urban scenes.
- **Uni**: single (uniform) orientation dominance. This includes vertical and horizontal dominance, but excludes axial dominance (as there are two orientations present). Can occur at any angle. Rather rare in regular scenes, but potentially useful for textures with diagonal dominance.

So far only the ori-kolumns have been used for a parametric description. Other maps are generated using the contrast and chromatic attributes; and combinations of maps are generated too. They are all written to a file with extension `.txm`, if flag `--saveTxm` is set.

Global Statistics and Blobs

A number of the texture maps is analyzed for its global statistics and the presence of individual regions. The global statistics appear as **Gst** in the saliency file and are useful as a triage in the structural analysis of a scene.

Blobs are regions in the texture maps. For a texture map, we find the connected components and determine its bounding box in the original image resolution. Since the texture map is of coarse resolution, ie. 32x32 pixels only, the corresponding bounding boxes in the original resolution are approximate only. The information for the blobs appears as **Blb** in the saliency file. Blobs often correspond to objects and scene parts.

1.3.5 Saliency

Saliency is the structural information that allows to decide what to analyse next. It is not a specific description, nor is there a particular algorithm that generates some saliency. It is rather a statistical summary of the outputted description which allows to choose what is of potential interest.

For example, if the image is full of texture, then we are interested in finding local variations by auto-correlating the kolumns. Conversely, if the image lacks any texture blobs, then any long contour in the pyramid of skeletons **ACNT** or bundles **ABNDG** is of potential interest.

On the other hand, if the image is full of texture, but contains a single, long contour at a lower level of the image space, then it is exactly that long contour that might be of interest. And in an image full of large shapes, the presence of an isolated blob might be the interesting structure.

The saliency file holds in particular information on texture and on shapes (Section 2.2.4). It does not hold any specific proposals from long contours or their groups, except of some attribute statistics. For a specific task, one would rather design an appropriate saliency calculating procedure. For example if we search for small objects, one would immediately focus on the texture analysis, see also Section 9.3.

1.3.6 Proposals

Proposals are a collection of descriptors that appear persistent across the image space. Presently, this is carried out for shapes and tetragons only (`shp` and `trg`). The descriptors and their bounding boxes are saved to files with extensions `qbbx` and `qdsc`, more details in Section 2.2.5.

1.3.7 Representation Formats

The descriptor attributes can be deployed in different forms. In the simplest form, we build histograms, suitable for traditional classifiers (LDA, SVM, RF). This allows fast classification that can serve to subselect candidates for more specific matching; in short, we build a cascaded recognition process. In the most elaborate form, we use the attributes in a multi-dimensional space, the attribute space as introduced above already. This space is suitable for identification of structure. We start with explaining the histograms.

Histogram of Attributes

One type of histogram was introduced already with the kolumns for texture analysis. Here we explain how the descriptor attributes are histogrammed more generally.

The attributes are histogrammed individually, with 5 to 12 bins each, both univariate and bivariate; a few trivariate histograms are generated as well, that appear under bivariate distributions. To exploit the position attribute, spatial histograms are formed as well, analogous to the texture analysis. Here they are formed for a small grid of non-overlapping cells, of dimensionality 3x3 by default. For each cell, both the univariate and bivariate distributions are formed. Thus, in total four types of histograms are generated:

flat, univariate	one-dimensional, taken for the entire image (or focus)
flat, bivariate	two-dimensional, taken for the entire image (or focus)
spatial, univariate	one-dim., taken from a 3x3 grid by default
spatial, bivariate	two-dim., taken from a 3x3 grid by default

They are also referred to as Histogram-of-Attributes. The total dimensionality is currently at ca. 24k for a 3x3 grid. Program `dscx` generates all four of them and they are saved in a separate file with extension `.hst`. Program `focsel` generates only the flat histograms and they are saved with extension `.hsf`.

The histograms from different images can be matched with the program `mhst` to quickly generate candidates for further, more detail matching (Chapter 4). The histograms can also be collected with program binary `collhimg`, suitable for training a traditional classifier. Those programs will be introduced later.

Kolumns In comparison, the kolumn histograms can be regarded as a fine version of spatial histogramming with a small set of attributes. Like the Histogram-of-Attributes, the kolumns can also be matched imagewise, in this case with a program binary named `mkoll`. Its use is analogous to the use of the program `mhst` and will be mentioned in Chapter 4.

Vector

The vector matching programs `mvec` and `mshp` exploit the attributes as a multi-dimensional space. The vector metric uses (in most cases) individual weights for the attributes, that is one weight per attribute, not for each descriptor. The metric includes the position and angle attribute by default, in which case, the representation can be called a *rigid vector template*. It is useful for identification of structure. Since structure sometimes appears at different orientation or with some variability, one can loosen the template by lowering the weight values, in which case the representation becomes rather a loose vector template. If we turn the weight values off for position (to value equal zero), then the representation is rather a statistical one, namely a set of structural elements. The matching programs allow to control those weight parameters. But since the metric favors equal list length, it is less suitable for allowing a 'controlled' subspace that would be useful for discriminating subtly different categories. For that it required the development of more flexible matching schemes. One approach will be introduced in Section 9.1.1.

1.4 Program Binaries

The program binaries are divided into computation, learning and utility binaries.

1.4.1 Computation

The following summarizes those binaries, that carry out computational processes. Many of them were mentioned already.

- **dscx**: carries out feature extraction and description as introduced already above and outputs various files that will be loaded by other recognition processes. Currently it takes only one image at a time and it is therefore somewhat slow due to repeated memory allocation for individual calls. Future versions will include processing a list of images for faster extraction.
- **mvec**: matches the descriptor vectors as outputted by **dscx**, and returns dissimilarity and similarity measurements for each pyramidal layer and descriptor type. It can be applied to any structure: shape, object or scene. There are two instantiations of vector-matching. One that matches only one pair of images (outputs of **dscx**), called **mvec1**. And one that matches one versus multiple, called **mvecL** ('L' for list).
- **mhst**: matches descriptor histograms and outputs a dissimilarity measurement. This can be used to identify candidates for vector matching (with **mvec**). This not only accelerates recognition, it most often improves prediction accuracy.
- **mkol**: matches kolumn histograms and operates analogous to **mhst**.
- **shpx**: carries out a refined segmentation for a target region using a color cue, and then describes the obtained region silhouette using the radial description (**rsg**), the arc description (**arc**) and straighter description (**str**).
- **mshp**: matches the description as produced by **shpx**.
- **motvec**: computes the motion vectors between descriptors of two frames (Section 3.3). This is essentially the same as **mvec1**, but outputs in particular the motion vectors between nearest neighbors. That allows to estimate motion flow.
- **knnv** [prototyped]: nearest-neighbor search of structures using a coarse-to-fine strategy to accelerate the retrieval process, as opposed to **mvec** above, that matches the entire pyramid between two structures (and is therefore slower). **knnv** starts with a matching of the top (pyramid) layers, then subselects images and progresses toward finer levels. Early experiments have shown that this strategy not only speeds up the search, but also improves the sorting. For the moment one can use the combination of **mhst** and **mvec** to accelerate recognition.

1.4.2 Learning

A learning process will be provided that determines category-characteristic descriptors by individual matching of vectors. In principle one can carry out such a learning procedure with the matching binaries **mvec**, but it is of course more convenient to have binaries that provide more automation.

- **kkcan** [planned]: searches for nearest neighbors across images of one category to find candidate descriptors. This is based on **mvec1** as introduced above.

- `kkgrp` [planned]: refines the search and selects a final set of category-characteristic descriptors using some clustering technique.
- `kkmtc` [planned]: matches the category-characteristic descriptors against a new image and outputs the degree of dis-/similarity per category-characteristic descriptor, similar to `mvecL`.

1.4.3 Utility

The utility programs rearrange the description output for specific purposes.

- `focsel`: selects descriptors from a region (focus), specified by the user as bounding box. Program `fochst1` extracts the histogram for one focus, program `fochstL` does so for multiple focii. Histograms are written to files with extension `hsf`. Program `focdsc1` extracts the descriptor spaces from one focus; they are placed into a file with extension `dsf`. The `dsf` or `hsf` files are loaded by matching programs such as `mvec`, `mhst` or `knnv`.
- `h2arr`: generates a single histogram array from the histogram file (as outputted by `dscx`), thus suitable for feeding directly to a traditional classifier (LDA, SVM, RF, etc.), introduced in Section 2.4.
- `collhimg`: same as binary `h2arr`, but taking a list of histogram files as input. It outputs a single matrix of size `[nImg nBins]`, suitable for training traditional classifiers.
- `d2vmx`: generates the vector matrices for the individual descriptor types to the format `[nDsc nAtt]`: number of descriptors times number of attributes. This matrix can then be deployed for clustering, ie. word formation. Introduced in Section 2.5.
- `collvec`: concatenates the vector matrices for a list of (image) descriptions to a single matrix of size, `[ntDsc nAtt]`, that is total number of descriptors (for all files) times number of attributes of a descriptor (Section 2.5).
- `ptchx` [prototyped]: Extracts rectangular patches of any size from one image using a list of specified bounding boxes. This serves to prepare the patches for shape extraction with `shpx`.

1.5 Demonstration, Varia

The following mentions the demonstration programs available, as well as administrative and utility routines for Matlab:

- `/DemoBaum` contains binaries that demonstrate the segmentation process, one for gray-scale analysis (as used in `dscx`), and one for a chromatic analysis using the RGB channels.
- `/DemoPlcRec`: contains Matlab scripts that carry out a simple place recognition experiment using programs `dscx`, `mvec` and `focsel`. It is an elaborate version of the example introduced in Section 1.2.2.
- `/DemoSgrRGB`: contains a binary called `sgrRGB` demonstrating the color-segregation process used for shape extraction (Chapter 8). It segregates an RGB image for a given target color, resulting in a black-white image with region boundaries that

are more precise than with the gray-scale information as used in program [dscx](#).
It also outputs the arcs and straighter descriptors for the foreground regions.
More information in Chapter 8.

- `/Demos`: contains various scripts demonstrating in particular parameter changes.
- `/UtilMb`: contains utility routines for Matlab.
- `/UtilPy`: contains utility routines for Python.

1.6 Recognition Pipeline

We have introduced a simplified pipeline already previously. With the full set of programs introduced above, we can now envision full pipelines for different tasks. We start with a generic recognition pipeline, and then suggest one for the task of text recognition in the wild.

In case of absence of a specific task, recognition may start with scene classification (*ScnClsf*), that allows to select candidates for scene identification (*Scnldfc*). Then one proceeds with a more local analysis, such as object or scene part recognition (*ObjClsf*, *Objldfc*), followed by a shape analysis:

```
dscx → ScnClsf(.hst) → Scnldfc(.dsc,mvecL) → fochst1(.slc) →
ObjClsf(.hsf) → Objldfc(.dsf,mvecL) → shpx → mshp
```

For the classification processes (*Clsf*) we have not named an algorithm, since there exists a large body of classifiers. If the goal is scene identification only, then one could deploy *mhst* to subselect candidates. If the goal is classification per se, then one can obtain decent results with the combination of PCA and LDA; a Random Forest classifier can sometimes achieve substantially better accuracies and is carried out much quicker during application. One could also try feature selection schemes.

For text recognition in the wild, the above generic pipeline is useful to generate approximate candidate locations. We here focus on the part of the pipeline, that attempts to find letter shapes. In that case it makes sense to classify individual shapes immediately, in order to locate character and text candidates as quickly as possible. For that we extract a focus at those shape locations (proposals), that are provided by the saliency file (*.slc*). We can deploy shape extraction and matching (*shpx* and *mshp*) to discriminate between letters and non-letters.

To achieve the highest identification accuracy possible, we apply the (demo) program *sgrRGB*, which returns the shape boundary, that in turn can be applied to an OCR network:

```
dscx → fochstL(.slc) → Clsf(.hsf) → shpx → mshp
→ sgrRGB(.slc) → OCR
```

Chapter 2

Descriptor Extraction [/DescExtr]

The program `dscx` outputs a number of files, some of which were mentioned already. They are written to subdirectory `/Desc` in our examples (of directory `/DescExtr`). The other subdirectories contain the following:

<code>/Desc</code>	output directory for description generated by <code>dscx</code>
<code>/Imgs</code>	sample images for immediate probing
<code>/Params</code>	example files for setting parameters
<code>/Regist</code>	text files containing lists of filenames
<code>/UtilMb</code>	Matlab routines and example scripts to read the outputted files
<code>/UtilPy</code>	Python scripts to read the output data files.
<code>/Vect</code>	output directory for vector files as generated by <code>d2vmx</code> et al.

The use of the program is explained first (Section 2.1), followed by explaining what type of data files it generates (Section 2.2). Then we introduce the available flags and options (Section 2.3). While the output is in a format that is suitable for our programs, it is less amenable for use with traditional classifiers (LDA, SVM, Random Forest, etc.). Sections 2.4 and 2.5 explain how the output files can be converted and collected to form vectors and matrices, that one then can directly feed to classification software, using the binaries `h2arr`, `d2vmx`, `collhimg`, `collvec`.

2.1 Program Use [dscx]

Two arguments are required, the image filepath and the output path for the data files:

```
> dscx pathImg pathOutFile
```

The input image can be `jpg` or `png`. The output path must include a slash (as the program checks for that). Here is an example,

```
> dscx Imgs/img1.jpg Desc/img1
```

in which the output file name `img1` is chosen to be the same as the image name, for convenience. This will then write the following files into directory `Desc`:

<code>img1.dsc</code>	descriptor attributes, used by <code>mvec</code> , converted with <code>d2vmx</code>
<code>img1.hst</code>	descriptor histograms, used by <code>mhst</code> , converted with <code>h2arr</code>
<code>img1.slc</code>	saliency information.
<code>img1.dsb</code>	descriptor bins, used by <code>fchst1/fchstL</code> .

The following files provide further description, but are written only if a flag is set:

<code>img1.kol</code>	histogram kolumns, used by <code>mkolL</code>
<code>img1.txm</code>	texture maps
<code>img1.qbbx</code>	bounding boxes of proposals
<code>img1.qdsc</code>	descriptors of proposals
<code>img1.dscRRE</code>	full set of ridge/river/edge attributes, used by <code>motvec</code>
<code>img1.dscCVP</code>	full set of arc/straighter attributes
<code>img1.BonPix</code>	boundary pixels
<code>img1.CntEpt</code>	the endpoints of ridge, river and edge segments.
<code>img1.Bbox</code>	bounding boxes of regions.
<code>img1.BonBbox</code>	bounding boxes concatenated across depth.
<code>img1.BonAsp</code>	aspects of boundaries

The last five files of the list - from `img1.BonPix` on - are introduced in Appendix A. The others will be explained in the upcoming sections.

The Matlab script `exsbDscxSimp.m` shows how to execute the program from Matlab, using fewest utility routines for clarity. The script `exsbDscxFull.m` is a demo including utility and plotting routines. An example for a wrapper routine is given with script `RennDscx.m` that also verifies proper termination of the program. The program output (`stdout`) should terminate with the string `EndOfProgram`.

Image Size Limits The program should work without failure for any image smaller equal 480x640 pixels. For larger images, our program assumes that the input image is a regular scene, and not an image made of excessive or artificial texture, white noise for example. The reason is that we operate with constant memory allocation for practicality. Since our process generates an enormous amount of features, we therefore limited memory allocation for larger images. For the CityScape collection with its 1024 x 2048 images, the program runs effortlessly, meaning no image has produced a feature output that comes close to the maximally allocated memory. The program should also work for images of size [3000 x 4000] such as an image taken by a cell phone. Yet it might fail with unusual image content; perhaps there exists a natural scene image (>480x640 pixels) with noise-like texture, that exceeds the allocation of contours. If you happen to encounter one, inform me on info@sehbau.com.

Other Limitations The program might fail if one choses unreasonable parameter values, as not all input parameters and arguments are verified for their proper interval.

2.2 Output Files

The directory `/DescExtr/UtilMb` contains Matlab function scripts to load the data files and to display the features. Example scripts are given in directory `/Examples` and deploy the output of the `dsc` file:

<code>exsbPlotDesc.m</code>	descriptor attributes, function scripts in <code>/Vect</code>
<code>exsbPlotHist.m</code>	histograms, function scripts in <code>/Hist</code>
<code>exsbPlotShape.m</code>	illustrates shape abstraction
<code>exsbPlotTtrg.m</code>	illustrates the tetragon-like shape abstraction
<code>exsbPlotBbox.m</code>	bounding boxes, routines in <code>/Bbox</code>
<code>exsbPlotBon.m</code>	region boundaries, functions in <code>/Bound</code>
<code>exsbPlotBonPix.m</code>	function scripts in <code>/Bound</code>
<code>exsbPlotSalc.m</code>	saliency information

The following sections explain how to load the individual files and will point out more example scripts.

2.2.1 Description Image (.dsc)

The description file with extension `.dsc` contains the descriptor attributes and can be loaded with routine `LoadDescImag.m`,

```
[DSC Kt] = LoadDescImag(pthVect);
```

Data structure `DSC` contains the following descriptor image spaces:

<code>.ACNT</code>	skeleton of contours (not full RRE set)
<code>.ARSG</code>	radial description of region boundaries
<code>.AARC</code>	selected set of arc segments (not full set)
<code>.ASTR</code>	selected set of straighter segments (not full set)
<code>.ASHP</code>	shapes (typically not touching borders)
<code>.ATTRG</code>	tetragons, preciser shape info of elongated shapes
<code>.ABNDG</code>	bundles, clusters of contours

They are loaded with a routine named `Read[Dsc]Spc.m`, that in turn calls a routine `Read[Dsc]Att.m`. For example for contours those routines are named `ReadCntSpc.m` and `ReadCntAtt.m`. The output can be manipulated as shown in script `exsbPlotDesc.m`. For the shape and tetragon descriptors there exist separate scripts, `exsbPlotShape.m` and `exsbPlotTtrg.m`. We proceed with explaining the organization of the attributes.

Attribute Organization

The attribute values are organized per type, not per vector (descriptor), more formally they are struct-of-arrays, not array-of-structs. For example, we extract the attributes of the first level of the space:

```
Cnt = DSC.ACNT{1}; % extracts the attributes for one level of the space
```

Then the attributes are available as fields, ie. the length attribute as an array in field `Cnt.Len`, the orientation attribute in field `Cnt.Ori`, etc. Thus, for the purpose of clustering or classification in Matlab (or Python), one has to concatenate them (horizontally) to a `[nDsc x nAtt]` matrix. The program `d2vmx` generates this matrix.

Subgroups of attributes are often organized as a 'pure' structure-of-arrays and loaded with the following routine:

```
[ARR szD] = ReadMtrxFlt( fid )
```

which loads the data structure as matrix **ARR** of size **[nDsc nAtt]**, namely number of descriptors times number of attributes.

Most attribute types are normalized to unit range, some only to approximate unit range in case of complex attribute definitions. Angle attributes come mostly in radians; the others normalized to unit range. Depending on the exact type of use of the vectors, one should certainly consider scaling them.

We firstly introduce attributes that are common to most descriptors. Then we introduce the details of the individual descriptor types.

General Attributes

The general attributes include the position of the descriptors in the image (or map), their chromatic values, their contrast, smoothness of the segment, size, angle, etc.

- **PosV**, **PosH** (also called **vpo**, **hpo**) : vertical and horizontal position. Typically the center of the feature, ie. the midpoints of contours and boundary segments; the pole for radial descriptors. Read with `ReadAttPos.m`. These values are normalized to $[0, 1]$, unlike the points provided below.
- **RGB**: chromatic red-green-blue triplet. Read with `ReadAttRgb.m`
- **Pts**: the key-points, such as the two endpoints and the point in-between. Read with `ReadDescPtsS.m` (short), or `ReadDescPtsF` (float). Usually in absolute (unscaled) coordinates, corresponding to the pyramid level.
- **Smo**, **Ter.Smo**: curve smoothness, for arcs and straighters. This is a local measure that expresses the proportion of smoothness in the curve. The measure is useful to discriminate between natural scene boundaries or boundaries of specular reflections versus boundaries of structure, which tends to be smooth. A L-feature is considered very smooth, because it contains only a small proportion of non-smoothness in its corner.
- **OrgCrv**: shape/region label. The shape from which the boundary segment was taken (arc or straighter).
- **OrgDth**: the depth map at which the boundary segment was obtained (0-indexing). Perhaps useful in an analysis progressing from high to low contrast.

Other attributes that are common to most (or some) descriptors are:

- **Les** arc length scaled, ie. for contours, arcs and straighters
- **Len** arc length absolute
- **Ori** orientation angle, $\in [-\frac{\pi}{2}, \frac{\pi}{2}]$ (half circle)
- **Dir** directional angle, $\in [0, 2\pi]$ (full circle)
- **Ctr** contrast, normalized or $\in [0, 255]$

In the following we survey the individual descriptor types.

Contour Attributes `ReadCntAtt.m`

The most valuable attributes are the length and angular orientation as mentioned above already. Another attribute is the degree of straightness `str`, which however shows little relevance for vector matching so far.

Rsg Attributes `ReadRsgAtt.m`

The most valuable attributes are the radius `rds`, the elongation `elo`, the degree of concavity `cncv`, and the circularity `cir`. The other attributes aid classification (with Histogram-of-Attributes from the `hst` file), but have not played a large role yet in vector matching.

Arc Attributes `ReadArcAtt.m`

The geometric parameters are loaded to variable `(.Geo)` as matrix `[nArc nAtt]`, with the above mentioned routine `ReadMtrxFlt.m`. The first column contains the curvature measure, which is the most valuable attribute. The others improve classification with histograms, but have not played a large role yet in vector matching.

Straighter Attributes `ReadArcAtt.m`

Contains the same attributes as for contours (above), but some other attributes are loaded as well, that are still under development.

Shape Attributes `ReadShpAtt.m`

Shape attributes are organized into several groups of attributes (loaded with `ReadMtrxFlt.m`). Two groups represent a description of the straighter segments of a shape, and are useful for scene analysis of indoor or outdoor scenes, where there exist often vertical and or horizontal structures. The first group is a coarse description using a 8-bin histogram of the straighter orientations, which then is analyzed for various axial alignments. The second group is a more refined description using a 12-bin histogram. In Matlab those two groups are read as matrices as follows, see function script `ReadShpAtt.m`:

```
[V.STR szL] = ReadMtrxFlt( fileID );
V.SFI      = ReadMtrxFlt( fileID );
```

where `.STR` and `.SFI` are coarse and fine information, resp. In the example script `exsbPlotShape`, those matrices are turned into data structures by the utility routine `u_AttsArrToStruct.m` and variables `AVEC.LabShpScors` and `AVEC.LabShpScors`, that contain the fieldnames:

```
Scors = u_AttsArrToStruct( SHPlv.STR, AVEC.LabShpScors );
Sfine = u_AttsArrToStruct( SHPlv.SFI, AVEC.LabShpSfine );
```

The data structure `Scors` has the following fields. Each one represents a degree of a specific structural bias, ranging from 0 (absent) to 1 (fully present).

Vrt	verticality
Hor	horizontality
Dg1	diagonal 1
Dg2	diagonal 2
Axi	axiality: both vertical and horizontal
Adg	axiality along diagonals
Vab	deviation from verticality
Hab	deviation from horizontality
Dab	deviation from diagonality
Tri	three axes dominating
Nil	no axis is dominating: shape irregular

This coarse information is useful for abstract classification. The fields of the finer description (data structure **Sfine**) are as follows. Some are the same in structural type as listed above, but not in value as we deal with a 12-bin histogram.

Vrt	verticality
Hor	horizontality
Vti	vertical with some inclination
Hti	horizontal with some inclination
Vob	vertical oblique
Hob	horizontal oblique
Dg2	diagonal 2
Dg1	diagonal 1
Axi	axial: both vertical and horizontal
Uni	one orientation
Dul	two dominant orientations (such as in axial)
Cvg	two dominant oris and converging
Agx	angle between the two most dominant orientations
Ori	orientation value of the (most) dominant orientation
Nil	no dominant orientation present
Dre	three dominant oris present
Vir	four dominant oris present
Fnf	five dominant oris present

Tetragon Attributes `ReadTtgAtt.m`

The geometric attributes are organized into four groups (as data structures of the output variable):

.GEOM	basic form parameters, such as length, elongation, etc.
.LAGE	alignment with respect to the vertical and horizontal image axes
.ANGS	angles for converging direction and intersection
.DICV	directional biases

A tetragon's cornerpoints and its axis are contained in the following variables:

- **.Cop** the four corner points aligned clockwise starting from upper left, followed by the upper right, etc. They are the intersection of the detected segments and therefore not necessarily congruent with the actual shape, for instance for a rectangle with round corners, the corner points would outline the rectangle as if it had sharp corners.

- **Ax** contains the two endpoints of the axis.

2.2.2 Histograms (.hst, .kol)

The Histogram-of-Attributes for an image is saved to a file with extension **hst**. It can be loaded with routine `LoadDescHist.m`, as demonstrated in example script `exsbDscxSimp.m`:

```
[HFU HFB HSP Nbin] = LoadDescHist([pthOut '.hst']);
```

where data structures **HFU**, **HFB** and **HSP** contain the histograms, flat univariate, flat bivariate and spatial (both uni- and bivariate), resp. Variable **Nbin** contains the number of bins used for each attribute, for the various types of histograms (**Uni**, **Biv** and **Spa**).

With program binary **h2arr** one can convert this output to a file that contains the histogram values as a single array (Section 2.4). The histograms of different images can be collected with program **collhimg** to create a matrix for training classifiers.

The dimensions for spatial histogramming can be changed through a parameter file. This will be explained under options for architecture (Section 2.3.1).

We recommend starting classification with the first two types, the flat univariate and bivariate histograms. The third type, spatial univariate, improves accuracy in particular for whole images. The last (fourth) type, spatial bivariate, did not consistently improve classification accuracy, when tested on various image collections. It is perhaps best to apply feature-selection schemes.

Kolumns Kolumns are saved to a file with extension **.kol**, if flag **--saveKol** is set. The kolumns can be loaded with script `LoadKolumns.m`. An example for loading and displaying them is given with script `exsbKolumns.m` in directory **/Demos**. Their use in scene categorization has not been tested yet; it probably makes most sense to concatenate them with the histogram of descriptors. In scene identification they provide decent accuracy for limited environments.

2.2.3 Texture Maps (.txm)

The texture maps for an image are saved to a file with extension **txm**, if flag **--saveTxm** is set. They can be loaded with routine `LoadTxtrMaps.m`, as demonstrated in example script `exsbTxtrMaps.m`.

```
TXM = LoadTxtrMaps([pthOut '.txm']);
```

The output variable **TXM** is a structure, that contains the following sets of maps:

- **KNT**: maps based on the count of segments per window.
- **OTX**: maps based on the dominant orientation present.
- **SAL**: saliency maps that are combinations of other maps.
- **CRM**: maps based on chromatic statistics.

Those sets of maps are organized as structures as follows:

Count Maps *KNT*

The data structure contains one map, *Num*, whose units correspond to the sum in a window. Another map, *Blk*, is a binary map whose ON-pixels signify the lack of segments, that is fewer than a minimum number of segments.

Orientation Maps *OTX*

Contains 7 maps that show dominance for the following orientation angles:

<i>.Vrt</i>	vertical orientation
<i>.Hor</i>	horizontal orientation
<i>.Dg1</i>	diagonal 1
<i>.Dg2</i>	diagonal 2
<i>.Axi</i>	axial: vertical and horizontal orientation (co-occurring)
<i>.Nil</i>	null: lack of a clear orientation
<i>.Uni</i>	one orientation is dominant (<i>vrt</i> , <i>hor</i> , <i>dg1</i> or <i>dg2</i>)

In regular scenes, the vertical, horizontal and null-orientation attributes are the most informative, followed by the axial attribute that shows in particular in urban scenes. In scenes or textures that appear at any orientation, the diagonal attributes and the uni-orientation attribute can be beneficial.

Saliency Maps *SAL*

These are maps that contain contrast information and combinations thereof, in particular with the orientation maps.

Chromatic Maps *CRM*

Some statistics of the red, green and blue channels.

As explained already in Section 1.3.4, the statistics and regions of some of these maps is summarized in the saliency file, as discussed next.

2.2.4 Saliency (.slc)

The saliency file can be loaded with routine `LoadDescSalc.m`; the example script `exsbPlotSalc.m` demonstrates how to read the variables. The loading routine returns four data structures:

```
[Txa Shp Ens Dsc] = LoadDescSalc( pthSalc );
```

that contain the following information:

- **Txa**: texture information based on contour statistics as explained in Section 1.3.4 and 2.2.3, which are particularly useful for detecting common scene textures and small objects.
- **Shp**: bounding boxes of salient regions, which are taken from the space of shapes in variable *ASHP*.
- **Ens**: an ensemble of proposals that is combination of the contour (**Txa**) and shape information (**Shp**).

- **Dsc**: statistical information of the descriptors, both the number of descriptors as well as some of their attributes.

We elaborate on what aspects each data structure holds.

Texture Data Structure **Txa**

The texture (data) structure holds information on the 7 biases, once for the entire map in structure **Txa.Gst** and once per blob in structure **Txa.Blb**. Then there exists also a data structure **Spt** holding dense point candidates. We recall that we have the following 7 texture biases:

```
aTxtBis = {'Num' 'Blk' 'Nil' 'Vrt' 'Hor' 'Axi' 'Uni'};
```

- **Gst** The data structure holds five statistical values for each texture bias (read with `ReadMapBisStat.m`): the proportion present in the map; and the minimum, maximum, mean and standard deviation. In a scene full of texture, the numerosity value **Num** is high for its proportion, its voidness **Blk** value is low. In a scene lacking contour texture, the inverse holds: low numerosity and high voidness. The example script `exsbSalBlobs.m` shows those two cases.

- **Blb** this data structure contains the blob information, the regions outlining contour texture. It is loaded with reading routine `ReadBlobOut.m` and contains the following structures and fields:

.Box	bounding box values in fields .Top , Bot , Lef , Rit .
.Typ	texture bias, numbering 1 ('Num') to 7 ('Uni'). 8: high contrast
.Cvg	coverage of image, $\in [0, 1]$

There may exist multiple blobs per texture type. The bounding boxes for **Typ=1** are the most general ones and outline any blob containing contour segments: if the boxes are small, than they outline small objects in isolation, demonstrated with example script `exsbSmlObjDet.m`.

- **Spt** contains a selection of points that represent clusters of high contour count and of high contrast, relative to their immediate context. They are plotted in the section called 'Spots' (in `exsbPlotSalc.m`), with routine `p_VisSearch.m`. These points are useful if the scene contains spots of moderate cluttering that we wish to locate, ie. when searching for objects.

Shape Data Structure **Shp**

This data structure holds the bounding boxes of selected shapes and their key aspects. They are read with reading routine `ReadShpOut.m`.

<code>.Box</code>	bounding box values in fields <code>.Top</code> , <code>Bot</code> , <code>Lef</code> , <code>Rit</code> .
<code>(.Typ</code>	irrelevant (NOT texture bias))
<code>.Cvg</code>	coverage of image, $\in [0, 1]$
<code>.Ctr</code>	contrast, $\in [0, 255]$
<code>.Cwd</code>	crowdedness, similar to texture bias ' <code>Num</code> '
<code>.Lev</code>	pyramid level from which a shape was taken, $\in [0, nLev - 1]$
<code>.IxShp</code>	index to shape for the given pyramid level (in <code>Lev</code>)
<code>.IxBon</code>	index to boundary (level in <code>Lev</code>)

Ensemble Data Structure `Ens`

The ensemble (data) structure holds a combination of the contour and shape proposals as well as an ordering of the their sizes:

<code>.Box</code>	bounding box values in fields <code>.Top</code> , <code>Bot</code> , <code>Lef</code> , <code>Rit</code> .
<code>.Typ</code>	descriptor type: 1-8 = contour texture; 10 = shape.
<code>.Cvg</code>	coverage of image, $\in [0, 1]$
<code>.Ctr</code>	contrast, $\in [0, 255]$. For contours set arbit. to 100. For shapes real value.
<code>.OrdGtoL</code>	order of indices from large to small (global-to-local)

Descriptor Data Structure `Dsc`

The statistics in this data structure are too innumerable to list them all. We here describe a few selected ones, that can be useful for fast orienting.

- `MxRngRR`, `MxRngEg` These two arrays contain the maximum contrast (range) value of ridges and rivers (RR) and edges (Eg) per level of the image space. This information is useful for autofocusing; if the RR values are extremely low (less than 10), then we likely face a blank visual field with perhaps specular reflections; or a scene with no clear contours, such as underwater scenes, or with regions of low contrast.

- `MaxSizSc1` contains the maximum size per level for each descriptor. If the one for shapes is large, ie. `max(MaxSizSc1.Shp)`, then we likely have a large object in the image center (if parameter `Shp.bordTouches = 0`, which is default); or it can be background, that is surrounded by objects, such that the background appears as an inside shape.

- `GryMmm` is a three-value array that contains the minimum, mean and maximum gray-level value for the entire image.

2.2.5 Proposals (`qbbx`, `qds`)

The proposals are saved to file if flag `--saveProp` is set, see example script `exsbProposals.m`. The loading routine returns a data structure `Bbx`

```
[BBx Nr] = LoadDescPropBbox( [pthOut Fixt.qbbx] );
```

that contains the bounding boxes selected from:

- `ShpGen` the entire shape space (`ASHP`).

- **TtgGen** the entire tetragon space (**ATTG**).
- **AxVrt** shapes that appear vertical (of **ASHP**).
- **AxHor** shapes that appear horizontal (of **ASHP**).

The attributes of the selected shape and tetragon descriptors are also extracted: they are saved to the **qdesc** file, which is loaded as follows:

```
[QDSC] = LoadDescPropAtts( [pthOut Fixt.qdesc] );
```

where structure **QDSC** contains the attributes as introduced previously; yet in this case they consist of a single data structure.

2.3 Options and Parameters

Options and parameters can be passed either as text file or as long options. For the use of a text file see the directory /Params. The file named `PrmDesc.Example.txt` contains the most important parameters. If one uses both, long options and text file, then the text file must appear as third argument, followed by long option arguments:

```
> dscx Imgs/img1.jpg Desc/img1 Params/PrmDesc.Example.txt --depth 4
```

whereby the values provided by the long options overwrite those provided by the text file. The naming of parameters is slightly different for the two types occasionally. On the long term, the naming as read from file will be the preferred one, as it is more systematic.

Many of the parameters will mainly regulate the number of descriptor vectors saved to the `dsc` file. For histograms, typically the full set of descriptors is taken (rre for contours, full for curve partitions), as the exact choice of those parameters has lesser influence on recognition accuracy. For manipulation with vectors however, some of these parameters can make a huge difference. In particular for place recognition I had observed substantial variations, but a full systematic search is still to be carried out.

In the following the parameter name for the text file is listed first, if existent already. If a long option is present, it is specified with a double dash `--` and is listed in parentheses (if the parameter can also be specified in the text file); single letter options are not in use.

2.3.1 Architecture

nLev (`--nLev`): number of levels of the image space. For a pyramid, the number is calculated automatically (with downsampling factor equal 2), whereby the top level does not subceed 16 pixels for one map side. For example, for a 256x256 image a five-level pyramid is generated: 256 (original resolution), 128, 64, 32 and 16. For a scale-space, the default equals five.

imgSpc (`--is`): sets the type of image space. By default it is the pyramid, with value equal one. The scale space can be selected by setting to value equal two. A demo script is in directory /demos called `exsbImgSpaces.m`.

HistSpaDim: sets the grid dimensionality for spatial histogramming. The default is 3x3. The dimensions are specified as rows and column, ie. changing to three horizontal cells that cover the entire image width, would be specified as:

```
HistSpaDim 1 3
```

Note that a grid of 5x5 generates a very high dimensionality with over 30k bins. The demo script `exsbSpaHist.m` gives an example (in directory /Demos).

Another parameter that can be considered part of the architecture is the window size for texture analysis. It will be mentioned further below, in Section 2.3.7.

2.3.2 Contours

Cnt.minCtr (**--cntMinCtr**): contrast threshold for contours. Default =0.05. This is the proportion, of the largest difference found in the range map of the gray-scale intensity image (for a 3x3 neighborhood).

The following two parameters - starting with **skl** - modify the output of the contour selection, the skeleton. To understand those changes see the example script `exsbContourSkel.m` in directory `/Demos`, or turn on plotting using flag **--plot**, which writes the image `ImgPyrSkel.png`.

Skl.MinSpc (**--sklMinSpc**): minimum spacing. Default =0.05. This parameter is specified as proportion of the image side length, ie. for an image side of 256 pixels it will be 13 pixels for the default value. Changes here can have a huge effect on performance, recognition accuracy in particular.

Skl.MinLen (**--sklMinLen**): minimum length. Default =0.05. This parameter is also specified as the proportion of the image side length. Changes here are less significant, in particular for large spacing values (set with **sklMinSpc**), as then only few short segments remain.

2.3.3 Regions

depth (**--depth**): depth of the segmentation process. Default **depth=3**. For **depth=1** no tree is grown: this corresponds to global thresholding only (with a single threshold). **depth=4** can be useful for large images, e.g. larger than 1000 pixel for one image side. A depth of five is the maximum.

Reg.minPixNode (**--regMinPixNode**): minimum number of pixels for a region to be segregated by the thresholding mechanism. This will affect the region count from the second segmentation map on. It will not affect the first segmentation map, as its regions are the first nodes of the tree.

Default equal 6. With larger values, processing occurs more rapidly, but may not capture fine texture properly anymore.

2.3.4 Radial Shape

--rsgMinPix: minimum number of boundary pixels for a radial region descriptor. The number is set for the original image resolution. For higher levels of the pyramid, a correspondingly lower number is used, namely **rsgMinPix-level**. E.g. for a value of 10, the higher pyramidal levels utilize values 9, 8, 7, ...

Default equal 4 for all levels. For larger values, ones risks loosing texture, thus if one interested in the global structure only, it can be useful to set higher values.

--rsgMinCtr: minimum boundary contrast for a radial region descriptor. The value is relative to the average contrast value for all extracted boundaries. Boundaries below that contrast value will not be described as radial descriptor.

2.3.5 Partitioning (Arcs/Straighters)

Parameter names starting with **cvp** regulate the partitioning process and therefore affect the outcome of both arc and straighter partitions. Parameter names starting with **arc** and **str** are specific to the respective descriptors.

--cvpMinSiz: minimum boundary size entering the boundary partitioning process. This is relative to the larger image side. For example a value of 0.02 for a [1024 x 2048] image will set the minimum size to 41 pixels. Default equal 0.05.

For small values this will seemingly regulate the number of straighter segments only, as those are rarer. For higher values, it will also start omitting arc segments.

Cvp.bInclBord (**--cvpInclBord**) [flag]: includes partitions at borders. The default is OFF, meaning that partitions at image borders are ignored. Turning them ON (by listing this option) can improve place recognition significantly.

Cvp.minCtr (**--cvpMinCtr**): minimum boundary contrast entering the boundary partitioning process. This is a proportion of the largest boundary contrast found in boundaries. Default equal 0.05.

Arcs and Straighters

Cvp.prpMinLenArc (**--arcMinLen**): minimum arc length to be described. This is relative to image side length. This is a better way of directly regulating the number of arcs than **cvpMinSiz**. Default equal 0.08.

--strMinGer: minimum straightness value for a straighter segment to be accepted. This is a fixed threshold $\in [0..1]$ based on the measure chord length divided by segment arc length. Default equal 0.8.

The following options set parameters for selecting the skeleton of partition segments, called here 'Gerust', abbreviated **gst/Gst**. Reducing the number of segments for finer scales often improves recognition - and also reduces matching duration. But I suspect that for smaller images (or higher levels of the pyramid; ie. smaller 100 pixels per side), that the selection might have less effect.

--arcGstMinSmo: minimum smoothness for arcs set for entire pyramid. Default approximately 0.20. This parameter is intended to eliminate segments (resulting from segregation of luminance gradients) that typically show high irregularity in their spatial course. Images of scenes taken at night tend to have those in particular.

--arcGstMinSpc: minimum spacing for arcs set for entire pyramid. Default approximately 0.05.

--arcGstOff: turns off any selection by setting all parameters values to zero (for smoothness, spacing and length). This then takes the full set of extracted arcs.

--strGstMinSpc: minimum spacing for straighters set for entire pyramid. Default approximately 0.05.

--strGstOff: turns off any selection by setting all values to zero. This then takes the full set of extracted straighters.

2.3.6 Shape (Arcs & Strs)

Shp.bordTouches: number of border touches permitted for a shape. Default equals zero, meaning only inside shapes are described. Values 1 to 4 permit the corresponding number of border touches for a shape.

(**Shp.minCtr**: minimum contrast for a shape. Not in use yet. Currently depends on the boundary contrasts measured.)

2.3.7 Texture

Txt.winSiz (**--txws**): side length of the rectangular window. The default is 16 and serves well for image sizes of 256x256 depicting regular scenes. For larger image sizes, an increase might make sense. For the search of small objects, a smaller size can be beneficial, see also the example `exsbSm10bjDet.m`.

2.3.8 Utility

--prms [flag]: displays the parameter values used. Default OFF.

--saveBbox [flag]: turns on saving of bounding boxes (`.Bbox` file) and contour endpoints (`.CntEpt`). Default OFF.

--saveBon [flag]: turns on saving of boundary information (files `.BonBbox`, `.BonAsp` and `.BonPix`). Default OFF.

--noBin [flag]: turns off saving of descriptor bins (file `.dsb`). Default ON.

--plot [flag]: plots contours and region boundaries for the entire pyramid; not available for scale space (image space `is=2`). Default OFF. The following images will be written:

- `Icnt.png`: contours plotted onto the color image for the original resolution.
- `ImgPyrBonOnly.png`: boundaries of the entire pyramid and for different depths.
- `ImgPyrCntOnly.png`: contours of the entire pyramid and for different levels.
- `ImgPyrSkel.png`: the selected contours (skeleton).

--verbose [flag]: for illustration or for tracking errors. Default OFF.

2.4 Collecting Histograms [*h2arr*, *collhimg*]

The histogram file can be converted to a text file, that contains the bin values as a single array, by applying program binary *h2arr* ('histogram to array'). The outputted array can then be used for classification. The program takes a *hst* file as input and an output file stem as second argument, for example:

```
> h2arr Desc/img1.hsf Vect/img1
```

This will write the file *img1.hari* to the directory */Vect*. It can be loaded as shown in *LoadHistImgArr.m*. A full code example is given with script *exsbH2arr.m*. The bin values represent the raw count; no scaling was carried out. The total dimensionality of the histogram can be over 24 thousand for 3x3 spatial histogramming. The example script *exsbSpaHist.m* in directoy */Demos* displays the histograms for different dimensionality.

With binary *collhimg* we can collect the histograms for multiple images. For *nImg* images, the program then generates a matrix of size, *nImg* x *nBin*, namely number of image histograms times number of bins. It can be deployed for training traditional classifiers such as LDA, SVM, Random Forests, etc. The same output can be achieved by calling *h2arr* individually and then concatenating the outputted histograms.

The program takes two arguments. The first argument is a text file containing a list of histogram files. The second argument is an output name for the matrix:

```
> collhimg ListHists.txt COLLHST
```

This will write the matrix into the file named *COLLHST.hstc*, where letter 'c' in the file extension stands for collection. The matrix is written to file in binary format (not as text). It can be loaded with routine *LoadCollHist.m*. An example script is in directory */MtchHst*, called *exsbCollHist.m*.

The second output variable of *LoadCollHist.m*, here called *Nbin*, contains the bin numbers of the individual histograms. The 4-element array *Nbin.Tot* holds the bin numbers for the four types of histograms, flat univariate, flat bivariate and the two spatial versions.

2.5 Generating Vector Files

Analogous to the pair of programs collecting histograms, *h2arr* and *collhimg*, there exists a pair of programs that convert the attribute values in the description file to vectors into matrix format, called *d2vmx* and *collvec*, respectively. Program binary *d2vmx* does so for one image (description to vector matrix). Program *collvec* does so for a list of images. The outputted matrix has the size number of descriptors (rows) times number of attributes (columns): *ntDsc* x *nAtt*, where the total number of descriptors *ntDsc* is for either one or several images. A corresponding label array is written as well, that contains level and image index.

These vector files are intended for manipulation of the description, ie. clustering, as discussed in Section 9.1.1, or for writing own distance measuring routines.

2.5.1 One Image Description [d2vmx] (.vecCnt, .vecRsg, ...)

Program binary `d2vmx` generates a vector file for each descriptor types. It takes as input a description file and requires the specification of an output file path, which in our examples is saved to the folder `/Vect`:

```
> d2vmx Desc/img1.dsc Vect/img1
```

This call then generates a separate vector file for each descriptor type with extension `vec[Dsc]`, ie. `img1.vecCnt`, `img1.vecRsg`, `img1.vecArc`, etc. The vectors are written rowwise, the columns therefore represent the attribute types. The output file path must contain a slash sign.

Level Array For each descriptor type, the program also generates a label array that contains the level index (of the image space) where the descriptor instance comes from. The indices are written with zero-indexing. The file extension is called `lev`, standing for 'vector level'. The full file extension are therefore `.levCnt`, `.levRsg`, etc.

The files can be loaded by routines `LoadDescVect.m` and `LoadDescVectLev.m`, respectively. An example script of how to run the program and read the output files is given in directory `/DescExtr`, called `exsbD2vmx.m`.

The script `o_AttsLabels.m` provides the attribute labels. Each descriptor contains two columns for the position information, `vpo` and `hpo`, for the vertical and horizontal position, respectively. The chromatic attributes are labeled `red`, `grn` and `blu`.

2.5.2 List of Image Descriptions [collvec]

This program allows to concatenate the vectors from multiple image descriptions into a single matrix. In simpler words, it is the list version for program `d2vmx`, but we call it `collvec` in analogy with program `collhimg` for histograms. The same result can be achieved by concatenating the output from `d2vmx`.

While `d2vmx` generates the vector file for each descriptor type automatically, here we run the program for *one* descriptor type only, specified with the second input argument. Thus, the program takes three arguments. The first argument is a text file containing a list of description (`dsc`) files. The second argument specifies the descriptor type, ie. 'skl', 'rsg', etc. The third argument is an output name for the matrix, called `COLL` here:

```
> collvec ListVec.txt skl COLL
```

This will write the matrix into the text file named `COLLVEC.skl.txt`, which can be loaded with script `LoadCollVec.m`. The program also writes a text file with labels to `COLLLab.skl.txt`, which can be loaded with `LoadCollVecLab.m`. The label array is of size `ntDsc x 3`, where the first column is the index of the level from which the descriptor was taken; the second column the image index; the third column is not used presently.

An example script is given in directory `/DescExtr`, called `exsbCollVec.m`.

Chapter 3

Matching Vectors [/MtchVec]

Program binary `mvec` matches the descriptor vectors of two images (`.dsc` files) as generated by the descriptor extraction program (`dscx`), or two focus files (`.dsf`) as generated by `focsel`. It can be used for any two structures expressed by the vectors, be it a scene, an object, a shape or a texture. The directories in the folder contain the following:

/Desc	description or focus files (as outputted by <code>dscx/focsel</code>)
/Imgs	sample images for immediate probing
/Mes	results of matching metrics
/Params	contains example files for setting parameters
/Regist	textfiles containing a list of filenames of description files
/UtilMb	Matlab routines

The program `mvec` matches the two list of descriptors using pairwise measurements and choosing the nearest neighbor. Two metric measures are available, a dissimilarity and a similarity value, abbreviated `dis` and `sim`, or sometimes also abbreviated as `dist` and `simi`, resp.:

- `dis (dist)` returns the Euclidean distance.
- `sim (simi)` returns the proportion of matches that are below a fixed threshold value, set with option `[dsc]TolMtc` or `tolMtc` for all descriptor types.

The program can be applied to two description files (`dsc`) as outputted by `dscx`, or to two description focus files (`dsf`) as outputted by `focsel`. The combination of a `dsc` and `dsf` file is not possible yet. The description files are required to have the same number of levels (pyramid height), otherwise the program returns no results (more flexibility to be included in the future).

The program comes in two variants:

- `mvec1` : matches one pair of images (or focii): one versus one. Its output is very elaborate, for example it generates nearest-neighbor information suitable for learning category-characteristic descriptors. The Matlab script `exsbFrames.m` demonstrates how to deploy the program and read its output. This program is useful for exploring parameter settings for different levels of the pyramid.
- `mvecL` : matches a list of images, or a list of focii: one versus multiple. Its use is demonstrated in Matlab script `exsbMvecLimg.m`. It outputs the measurements

per descriptor only and not for the entire pyramid (as in `mvec1`). It is useful for matching at large scale.

Then there exists also the binary `motvec`, that calculates the motion vectors between nearest neighbors. It is similar to `mvec1`, but returns the vectors only, useful for estimating motion flow, introduced in Section 3.3.

3.1 Program Use [`mvec1`, `mvecL`]

We firstly explain the use of `mvec1`, the matching of two description files. Their file paths are given as arguments. For example for two description files from images (`.dsc` from `dscx`) we write:

```
> mvec1 Desc/img1.dsc Desc/img2.dsc
```

Or for two focus file (`.dsf` from `focdsc1`):

```
> mvec1 Desc/foc1.dsf Desc/foc2.dsf
```

The description files are required to have the same number of levels, ie. generated by similarly sized images, otherwise it returns no results. The output will be further explained in Section 3.2.

For the use of the program `mvecL`, we specify a file path as well as a text file that contains the file paths for a list of files to be matched with. We prefer to keep those text files in a folder called `/Regist` (for register):

```
> mvecL Desc/img1.dsc Regist/FinasImg.txt
```

This will write the metric measurements into a file named `Vec.txt` in directory `/Mes`. One can specify a different file name by giving a third argument, ie.:

```
> mvecL Desc/img1.dsc Regist/FinasImg.txt Mes/ImgVec.txt
```

By default, the program matches all descriptor types for the entire pyramid. The options allow to subselect descriptors and levels, as well as to set attribute weight values.

3.1.1 Options

Options can be set by file or by long options (as in case of `dscx`). Again, the use of a file is explained through the Matlab scripts. In the following the use of the long options is explained.

The first list of options set a parameter value to the same value for all descriptor types (Section 3.1.1), which can be unspecific in some cases, but which is convenient for coarse tuning. The second list of options allows to adjust parameters of individual descriptor types for fine tuning (Section 3.1.1).

General (All Descriptor Types)

The following options set values for all descriptor types simultaneously and are useful for coarse tuning. Default values are given in approximate values only, as they are often individual to the descriptor type:

--tolMtc [similarity metric]: sets matching tolerance to a fixed value, arccos all pyramidal levels (and across all descriptor types). This is for the similarity metric only (section `Simi` in output). Default: ca. 0.05.

--wgtRGB: sets the weight value for the RGB difference. Set this parameter to zero if chromatic information is irrelevant, for example when places are to be recognized at either day or night. Keep in mind that three difference values are taken (R,G,B) and that this weight parameter therefore has more influence than the others. Default: ca. 1.0.

--wgtPos: sets the weight value for the position parameter for each descriptor type. A value of 0 turns off the influence. Default equal ca. 1.0.

Individual (Per Descriptor Type)

--cntTolMtc: tolerance for contour matches, for the similarity metric. Fixed value for all levels, but will try to provide something more flexible. Default: complicated.

--rsgTolMtc: tolerance for radial descriptor matches. Analogous to option `cntTolMtc`.

--arcTolMtc: tolerance for arc segment matches (see `cntTolMtc`).

--strTolMtc: tolerance for straighter segment matches (see `cntTolMtc`).

More in progress.

Utility

--prms: displays the parameter values used.

3.2 Output**3.2.1 Program `mvec1`**

Two types of results are returned. One type are the measurements of the list-matching metrics, written to `stdout`. The second type are the nearest neighbor indices (Section 3.2.1). See the example `exsbFrames.m` for the upcoming explanations.

Descriptor List

The list-matching measurements (appearing in `stdout`) are returned once for the entire pyramid, that is for each level (and each descriptor); and once integrated, for each descriptor type; as well as for the total, called image measure. The values per descriptor type and per image look as follows:

```

----- desctypes -----
dty  dis      sim
skl  1.357370  0.117928
rsg  1.480588  0.098934
arc  1.489363  0.009344
str  1.351547  0.054755
shp  5.827399  0.048974
eodty.
----- img -----
dis 23.574306
sim 0.000000
eoim.

```

The strings `----- desctypes -----` and `----- img -----` help locating the beginning of the respective measurements sets; as well as the 'end-of' strings `eodty` and `eoim`. This is carried out with the Matlab script `u_MtrMesSecs`. The actual measurement values are read by routine `u_MtrMesScnf`.

Nearest Neighbors

The nearest-neighbor indices are written to files in directory `/Mes` with prefix `NNspc` and suffix `12`, the latter indicating direction of comparison. The indices are kept separate for each descriptor type, ie. `NNspcCnt12`, `NNspcShp12`, etc. They are loaded with scripts `LoadNNDspace` and `ReadDescNNs`. The section 'Correspondence' in script `exsbFrames.m` establishes a visual correspondence for illustration.

3.2.2 Program `mvecL`

The results are written to three files into directory `/Mes`, called `Vec.txt`, `MesDtyDis.txt` and `MesDtySim.txt`. Each row corresponds to one pair of matching.

- `Vec.txt` contains the metric measures for the ensemble of descriptor types. It consists of four columns, where the first is the dissimilarity value, the second the similarity value. The third and fourth column are empty (zero) for the moment. The Matlab routine `LoadMtchMes.m` shows how to load the file.
- `MesDtyDis.txt` contains the dissimilarity value for each descriptor type. Each one consists of 7 columns, corresponding to contour, radial shape, arc, straighter, shape, tetragon and bundles of contours.
- `MesDtySim.txt` contains the similarity value for each descriptor. Its organization is the same as for dissimilarity values.

With the output of files `MesDtyDis.txt` and `MesDtySim.txt` one can develop individual ensemble measures.

3.3 Motion Vectors [motvec]

The program `motvec` takes as input two description files and the arguments are therefore specified analogous to program `mvec1`:

```
> motvec Desc/frm1.dsc Desc/frm2.dsc
```

This will output the vectors into directory `/Mes` in file `A.motvec`, which can be read as demonstrated with script `LoadMotVec.m`. By default it will use the skeleton contours as provided in those description files (structure `ACNT`). The program will also search for the presence of the `dscRRE` files in the same directory (in our example `/Desc`); if they are present, then the RRE vectors are taken for motion calculation, and the skeleton descriptors are ignored.

Since the program `dscx` does not save the RRE set by default, this has to be triggered by the corresponding flag. Example script `exsbMotVec.m` demonstrates how to do that, ie. by setting `OptK.saveRRE = 1;`.

Chapter 4

Matching Histograms [/MtchHst]

There are two program binaries matching histograms. One for the Histogram-of-Attributes, and one for the texture kolumns, called `mhstL` and `mkoll`, respectively. Their use is analogous and we therefore firstly explain the usage of the former, and then mention the latter (Sections 4.1 and 4.2, resp). The use of both is exemplified with the script `exsbMhstPlcRec.m`.

4.1 Histogram-of-Attributes [`mhstL`]

The program `mhstL` matches an attribute histogram against a list of other attribute histograms. This coarse comparison allows to identify a selection of candidates, that are then applied to a classifier, or that are immediately used for more accurate matching with `mvec`.

The program takes as input a histogram file (`.hst` file) as generated by the descriptor extraction program `dscx`, or as generated by focus selection program, `fchst1` (or `fchstL`). The directories in the folder contain the following:

<code>/Desc</code>	histogram files (as outputted by <code>dscx/focsel</code>)
<code>/Mes</code>	results of distance measurements
<code>/Regist</code>	textfiles containing a list of filenames of description files

The program `mhstL` calculates the Hamming distance between histograms and outputs the array of measurements to a file. Its input arguments are analogous to those of the program for vector matching (`mvecL`).

4.1.1 Program Use

Analogously to the program `mvecL`, we specify a histogram file (`.hst`) and a text file containing the file paths for a list of histogram files to be matched with:

```
> mhstL Desc/img1.hst Regist/FinasImg.txt
```

This will write the metric measurements into a file named `HstLst.txt` into directory `/Mes`. One can specify a different file name by giving a third argument, ie.:

```
> mhstL Desc/img1.hst Regist/FinasImg.txt Mes/MtcImg.txt
```

The histogram files need to originate from images of (almost) similar size, in order to properly match the spatial histograms.

By default, the program matches all descriptor types, ie. contours, arc segments, straighter segments, etc. as that typically yields best results. Matlab script `exsbMhstL.m` (in the main folder), gives an example of how to run the process.

The program can also take focus histograms (`.hsf`) as input:

```
> mhstL Desc/img1_f1.hsf FinasFoc.txt
```

in which case the list of filenames in `FinasFoc.txt` need to be `.hsf` files as well. For the focus histograms their original size is irrelevant, as no spatial histograms are involved; the idea of focus histograms is to create individually sized spatial histograms. Thus, the user must determine, whether it makes sense to compare two focii of different size.

4.1.2 Output

The measurement values are written twice to file, once ordered and once unordered, both in text format. The ordered output, written to `Mes/Hst.txt`, contains two columns, one containing the index of order in zero-indexing format, and the other the distance values. The ordered file, named `Mes/HstUor.txt`, contains the measurement values in unsorted order without any other information.

4.2 Kolumns [mkoll]

As mentioned already, the usage of this binary is analogous to the use of `mhst`. We specify a kolumn file (`.kol`) and a text file containing the file paths for a list of kolumn files to be matched with:

```
> mkoll Desc/img1.kol Regist/FinasKol.txt
```

This will write the metric measurements into a file named `KollSt.txt` into directory `/Mes`.

Kolumns are generated for the entire image only. There is no focus selection. An example of how to use this binary is given with script `exsbMhstPlcRec.m`.

Chapter 5

Focus Selection [/FocSel]

The goal of focus selection and matching is to accelerate visual browsing by maximally exploiting the descriptor output, see again the recognition pipelines suggested in Section 1.6. Structure appearing outside the image center can always be more precisely described and recognized, if we perform a saccade (change of camera direction) and then launch the descriptor extraction process again; or if we use the shape-extraction program `shpx` (Chapter 6). But that also includes more computation without having made full use of the feature description (from `dscx`). To fully exploit the richness of the descriptor output, the process of focus selection subselects the descriptors of a specified region, the *focus*, and saves them to a file for further analysis, ie. to be matched with `mvec` or `mhst`.

The program takes as input a description file (`dsc`) and a region specified as bounding box. The program then identifies the descriptors contained in that bounding box, and saves them to the file, the so-called focus file. The selection process comes in three variants:

- `focdsc1` : extracts the subset of the descriptor spaces for one focus, and saves them to a file with extension `dsf`. That file can be loaded by `mvec` for full vector-by-vector matching.
- `fochst1` : generates an attribute histogram for one focus, and saves it to a file with extension `hsf`, useful for rapid classification.
- `fochstL` : generates attribute histograms for a list of focii as specified in a text file.

The histogram output of the latter two programs allow for rapid classification and therefore rapid visual orienting. The Matlab scripts `exsbFocDsc1.m`, `exsbFocHst1.m`, `exsbFocHstL.m` and `exsbFocDscFew.m` demonstrate how to deploy those programs. They call the corresponding wrapper functions named `RennFocDsc1.m`, `RennFocHst1.m` and `RennFocHstL.m`

The folder contains the following directories:

<code>/Desc</code>	description files (as generated by <code>dscx</code>)
<code>/Focii</code>	output directory for focal selections, the <code>.dsf</code> files
<code>/UtilMb</code>	Matlab scripts to read the output data files

The process extracts the corresponding part of the original feature pyramid. To illustrate that, we refer to the scheme in Section 1.3.2. For example selecting from a

quarter region of the image (with four levels), the corresponding subset of the list is organized as follows,

```
lev 2      ||
lev 1      ||||
lev 0      |||||
```

and starts with level 0. It reaches only level 2, as we extract a subset of the pyramid. The extracted focus pyramid has therefore three levels, `nLevFoc=3`.

5.1 Program Use [*focdsc1*]

Firstly the use of *focdsc1* is explained. It takes three arguments:

- 1) a description file as generated by *dscx*. The file name *must* include extension *dsc*.
- 2) bounding box parameters specified as *top bottom left right*.
- 3) [optional] an output file name *without* extension, where the selected descriptors are written to.

Example: We extract a 40x40 region (height x width) from the upper left:

```
> focdsc1 Desc/img1.dsc 10 50 10 50 Focii/foc1
```

The selected descriptors are then saved to directory */Focii* appending the extension *dsf*. If the output file name is not specified, the program will write to the file called *Focus.dsf* in the same directory. How the output file is loaded will be explained in the next section.

The program calculates the number of pyramid levels automatically for the selected region. If no descriptors are found in that subspace, then no output file will be saved. More details on its output will follow below.

The program *fochst1* generates a single histogram, which is written to a file with extension *hsf1*. The program requires the bin file with extension *dsb* to be present (in the same directory as the *dsc* file), as generated by *dscx*.

The program *fochstL* does the same as *fochst1* but for a list of *focii* specified in a text file, named *BboxFocii.txt*, ie.

```
> fochstL Desc/img1.dsc BboxFocii.txt FOCII1
```

in which the bounding boxes are given rowwise. The first line of the file must contain the number of bounding boxes to be read. The output is written to file *FOCII1* in the example.

5.2 Output

The program writes both standard output and file output.

The standard output contains the number of levels, that was calculated automatically, `nLevFoc`, and the total number of descriptors detected in the focus, `ntDsc`, for example:

```
nLevFoc 3  
ntDsc 27
```

If no descriptors are found, the output returns `ntDsc 0` and no data are written to file.

The selected description is written to the `dsf` file. As mentioned already, if no output name is specified, the values will be written to a file named `Focus.dsf` in the same directory. The `dsf` is similar to the `dsc` file as generated by `dscx`. It is loaded as demonstrated in script `LoadFocDesc.m`, located in directory `/UtilMb`.

The example script `exsbFoc1.m` plots both the description of the entire image, and those of the focus (selection). The plotting routines (`PlotCntSpc.m`, `PlotRsgSpc.m`, ...) are found in directory `/Plot` of the folder for descriptor extraction (`/DescExtr`).

Chapter 6

Shape Extraction [/*ShpExtr*]

The goal of shape extraction is to obtain a more accurate description for a certain shape silhouette, than it was obtained with *dscx*. To achieve this, we specify the RGB triplet of a region (silhouette) under investigation, for example taken from the shape descriptor *shp* of the description file. That region can represent anything of interest, a simple shape itself, an object silhouette, an object part, or scene part.

Using the specified RGB triplet, program *shpx* then carries out a simple color-segregation process, whose output can be studied with the demo program *sgrRGB*, to be introduced in Chapter 8. This process achieves a finer segmentation than that was obtained with *dscx*. Program *shpx* then describes the shape as was done in *dscx*, namely by descriptors *rsg*, *arc* and *str*, whose attribute values will be more accurate due to the finer outline. In addition to those descriptors, *shpx* also returns the spectra of the local-global space. The entire description is then saved to a file with extension *shp*, that can be used with the shape-matching program *mshp1* (next chapter).

6.1 Program Use [*shpx*]

The input to the program is an image or an image patch. Taking an entire image may make sense, when the shape is large and covers almost the entire image width or height. If the shape under investigation is smaller, than it is of advantage to crop the image to its expected size, ie. extracting the individual letters of some text (in the wild). Cropping helps eliminating any distracting region. Cropping should not be too tight, as that makes proper curve description difficult; a border of at least 1 pixel should be included, even if we know the exact shape outline. By specifying an RGB triplet, we ensure that we deal with only that one shape. Thus, we specify for example:

```
> shpx Imgs/ptch1.jpg 150 100 185 Desc/ptch1
```

where *Imgs/ptch1.jpg* is the image patch; 150, 100, 185 is the color triplet; and *Desc/ptch1* is the output file stem. This will write the shape file *ptch1.shp* to directory */Desc*. The example script *exsbShpx.m* runs two patches of letter shapes. If a parameter file is provided, it must follow the output file, before any long options are specified. The script rather ensures proper operation of the program; a more illustrative example is given with the demo program *sgrRGB*.

Depending on the complexity of the image patch, it may well be that the segmentation output identifies other, smaller regions of similar color in the image patch than

the one aimed at. The matching program will take the largest one for matching.

6.1.1 Parameters and Options

Setting parameters and options is analogous to program `dscx`. An example of how to set parameter values is given with the file `Params/PrmShpx.Example.txt`. The parameter filename must contain the string `PrmShpx`. Flag `--prms` displays the parameter values used (default OFF); flag `--prmchg` displays the parameter values read by the parameter file (if provided).

6.2 Output

The shape file with extension `shp` is the only file output. For other types of output, one can deploy program `sgrRGB` (Chapter 8).

Chapter 7

Shape Matching [/ShpMtch]

The shape matching program `mshp` correlates shape descriptions as generated with program `shpx` (as introduced in the previous chapter). For a pair of shapes, three groups of measures are calculated.

- 1) descriptor distances: arc, straighter and radial-shape. The metric includes the angle and position attribute by default, but can be turned off by providing a parameter file containing preferred weight values.
- 2) spectral differences: one for the bowness spectrum and one for the straighter spectrum. This information is independent of orientation and position.
- 3) ratio of sizes: perimeter, height of bounding box, width of bounding box.

The measures will be outputted separately and the user can combine them to an ensemble as desired.

For the moment there exists only the one-on-one version, `mshp1`, whose usage is analogous to the vector-matching program `mvec1`.

7.1 Program Use [mshp]

Analogous to the vector matching program `mvec1`, we specify two files, in this case shape files with extension `shp`:

```
> mshp1 Desc/A.shp Desc/B.shp
```

A parameter file can be specified, that must contain the string `PrmMshp`. The example script `exsbMshp1.m` gives a simple demonstration of how to provide the parameter file and how to read the measures from the standard output.

7.2 Output

There exists only standard output. It displays the three groups of measures in three separate lines, ie.

```
mes 1.681 4.189 0.008
rts 0.847 1.033 0.969
spk 0.104 0.104
```

In detail, the values are as follows:

- **mes**: the three values arc distance, straighter distance, radial-shape distance. They can be combined to an ensemble as desired. Multiplication is a safe and uncomplicated way to obtain a reasonable prediction accuracy. A weighted sum may perform better, if the appropriate weights can be found.
- **rts**: the three values correspond to the ratios of the perimeter, height and width.
- **spk**: the two values correspond to difference of the bowness and straighter spectrum.

Chapter 8

Demo Segregation RGB

[/DemoSgrRGB]

The program `sgrRGB` segregates a color image into two groups according to a specified RGB triplet. Its purpose is to obtain a more precise outline of an object or scene part, than obtained by the region segmentation process (in `dscx`) that occurred in gray only. For example, we have detected a specific shape in the descriptor output, ie. in `rsg`, `shp`, `ttg`, etc. Since we know the approximate color of that shape, we can use it to cue the color-segregation process, and we obtain a more precise shape outline. The target color is assumed to represent the foreground; then, the segregation process automatically calculates a distractor target, assumed to be background. The process then performs a 2-Means clustering with a single pass (iteration). The segmented shapes of the foreground regions are then partitioned and described as it is done in program `dscx`, but only for the original resolution (no pyramid is generated, no divisive segmentation takes place).

To improve and accelerate the segregation process, the image should be cropped to the approximate size of the target object (or scene part). That is, the process is rather applied to only part of an image, called *patch* hereafter. The program will determine a distractor triplet using the image border pixels and for that there are different initializations possible.

The example script `exsbSgrPtch1.m` demonstrates how to run the program and how to load the output files. The script `exsbSgrPtchInit.m` compares the output for the different initialization techniques.

8.1 Program Use [`sgrRGB`]

The program `sgrRGB` takes a (color) patch as input and a RGB triplet, ie. {150, 100, 185}:

```
> sgrRGB imgs/ptchX.jpg 150 100 185
```

This will generate the following files (there is no option for specifying a filename yet):

<code>BonFore.bonPix</code>	boundary pixels for the target/foreground
<code>CrvPrt.cvp</code>	curvature partitions (arcs and straighter segments)
<code>Mlab.mpu</code>	black-white map, with white the foreground
<code>Ifore.png</code>	the foreground regions in average color
<code>BonBack.bonPix</code>	boundary pixels for the distractor/background

How those are loaded will be explained in Section 8.2. In the following it is explained what options are available.

8.1.1 Options

The following long option is available, to be specified with a double dash '--'; single letter option are not in use.

--init: the initialization technique that calculates the distractor triplet, specified as digit 0, 1 or 2:

- 0 (default): uses the (image) patch borders to calculate the average RGB value. The complexity corresponds to the number of border pixels: $O(nPxBorder)$.
- 1: determines the farthest distance between target and each border pixel. The complexity is the same as for `init=0`: $O(nPxBorder)$.
- 2: measures the distance of the target triplet to black $\{0,0,0\}$ and white $\{255,255,255\}$, and takes the farther of either as the distractor. This has lowest complexity as it calculates only two distances: $O(2)$.

For small patches with homogeneous regions, such as a character shape in text, there is little difference between the different types of initializations. The larger the patch and the more heterogeneous the colors, the greater the differences. Since the example script runs the program on an entire image, one can observe substantial differences between different initializations.

8.2 Output

The `.cvp` file is loaded by routine `LoadCrvPrt.m` in `/DescExtr/UtilMb/Curve`.

The `.mpu` file holds the size of the map in the first two integer values, height and width. The remaining values are of type `unsigned char` and describe the map values. It is loaded with Matlab script `LoadMapUch.m` (in `/SEHBAU/UtilMb/FileIO/`). The `BonPix` files can be loaded as explained in the appendix (Section A).

Chapter 9

Applications

We have already given an example of a generic processing pipeline for scene recognition, and one for text recognition in the wild (Section 1.6). Here we mention further applications, which will also clarify the strengths and short-comings of the system implemented so far. We firstly discuss two prominent tasks, namely categorization and identification (Section 9.1). Then we discuss the use of structural description for navigation (Section 9.2). Finally, we highlight some other applications (Section 9.3), as well as possibilities of a fusion with other methodologies (Section 9.4).

9.1 Classification and Identification

For the classification or identification of an arbitrary scene, the pipeline as mentioned in the introduction (Section 1.6), is the most straightforward way to implement either process. Firstly we classify using the `hst` file. As pointed out already, the combination of PCA and LDA provides often good results for abstract categories, often in the range of Deep Networks that were optimized for resources (and typically have lower recognition accuracies than full Deep Networks).

The histograms that we provide can be considered generic. For specific categories, different bin counts and edges for attributes can make quite some difference, as well as the choice of generated bivariate histograms and the dimensions of spatial histogramming. To explore such parameter variations we offer only a few options so far. But one can easily develop one's own histogram variations, using the vector matrices as outputted by programs `collvec` and `d2vmx`.

Since many of our attributes describe structural biases of almost binary character, in particular the shape attributes, it is worthwhile trying classifiers that perform for such data better, such as Random Forests or Deep Belief Networks, in particular the Restricted Boltzman Machine (RBM). But also the careful development of an ensemble classifier has quite some potential to improve the classification accuracy.

With the obtained classification posteriors, one can subselect candidates for matching vectors (using `mvec`). Although the vector-matching process is better suited for identification, one can try to improve the classification result by a Nearest-Neighbor analysis.

For subtly different categories, such as the Indoor-67 collection, it requires the development of representations, that exploit the vector space in full, something that requires careful development, ie. using `mvec1` or clustering (upcoming below). A

pragmatic choice would be to combine the result with a Convolutional Neural Network (CNN), that effortlessly learns subordinate categories (albeit at high cost and lack of structural interpretation). For that to be successful one would have to train the CNNs according to the confusions of the classification results.

For the identification of a scene, matching vectors is the immediate choice ([mvec](#)) and can be improved by matching individual regions (focii, shapes). This will work well in indoor scenes of a limited environment, for example a house or the floor of a building.

For the identification of scenes from any environment, a quasi-world recognition, the methodology of Local Features is still better. It is worthwhile combining the two approaches (Section 9.4.1). For example, one could take histograms of gradients at specific locations in the image, such as the regions provided in the saliency file ([.slc](#)), or of the contour skeleton ([ACNT](#)).

9.1.1 Clustering

The collected vectors (Section 2.5) can be deployed to build representations that lie between the abstract histogram representations and the rigid vector template. Analogous to 'words' in the Local Feature approach, one can determine characteristic vectors that then are utilized to express categories. In the Local Feature approach these words are typically determined from the entire training set, spanning all categories. That generic word formation is not specific enough in this case, as the attribute space of a descriptor type is of rather small dimensionality in comparison to the high-dimensional histogram of gradients. Instead, this approach requires the clustering per category as we look for category-specific structures. The procedure is as follows.

We deploy program [collvec](#) to build a attribute matrix for each category. Then we search for densities in the attribute space, which constitute category-characteristic vectors. The number of those vectors naturally varies between categories. When we apply those vectors to a testing sample, then we determine the number of best matches for each category and choose the category with the most matches.

9.2 Recognition for Navigation

For an efficient navigation through an environment, it is beneficial to detect vertical and converging structures. Contour information offers a fast access to distant and thin structures, region information often holds converging lines better than contours.

Vertical structure In road scenes, there exist often thin structures at distance, ie. lamp posts, traffic sign post, poles, etc. Those appear as ridge, river or edge contours, whereby ridge and river contours conveniently delineate a structure's axis. A fast way to obtain hypotheses is to deploy the contour skeleton [ACNT](#) of the vector file and integrate across the pyramid to find strong hypotheses. Since the skeleton is already a reduced contour set, there exists the risk of misses. A more thorough way would be to use the full RRE set, which can be loaded from the [dscRRE](#) file.

In indoor scenes, structures are nearer and have larger width, ie. walls, door frames, furniture, etc. In that case, we firstly focus on region boundaries and their descriptions, such as radial descriptors [ARSG](#), shape descriptors [ASHP](#) and tetragon descriptors [ATTG](#). In particular the latter gives us an immediate understanding of the

geometry of our surround; many of the attributes were designed for that, many more could be developed.

Converging structure Converging lines often exist as boundary segments of regions, that represent scene parts. In indoor scenes, those segments are often part of the shapes as described in list [ASHP](#). As explained, [ASHP](#) contains by default only the inside shapes, ie. those not touching any image borders. When no inside shapes are present, as is the case in a wide-open road scene, then we include those touching the border by increasing the parameter value [bordTouches](#). Should this still not provide candidate segments, we then analyse the full set of straighter segments, loaded with routine `LoadCVPfull.m` (turn ON saving for full set with long option `--saveCVP`).

Road Surface The region segmentation process easily returns the regions corresponding to the pavement. The segmentation process is so sensitive, that it will also return any pavement patch with slightly different color, ie. originating from road repairs or from frequent traffic such as the rills from truck wheels. The latter is often described by ridge or river contours as well.

The challenge with this sensitive output is to distinguish between task-relevant and task-irrelevant regions. For an analysis of road surface with respect to potential damage, any region of the surface might be relevant; for driving, we want to discriminate between non-obstructive patches and those, that pose a potential danger such as a slippery surface, spilled oil, etc. Since the descriptors are generated for higher contrasts by default, it is likely that some low-contrast regions are not described and therefore not available as radial shape [rsg](#) or abstracted shape [shp](#). And since those descriptors probably are not sufficient for an exhaustive characterization, it is best to analyze the whole set of boundaries, available in file [BonPix](#) (Section A.3).

If the scene is a dirt road, whose tracks (from previous vehicles) appear nearly equal to its surround, and therefore lacking border candidates from regions, then it is useful to focus on ridge and river contours, ie. using the full set in the [dscRRE](#) file. In this case, preprocessing the image might be of benefit, such as smoothing (low-pass filtering) the image; this facilitates finding contiguous ridge and river contours, albeit at the risk of obtaining contours that connect too much structure - as is the case for edge detection for different scales.

9.3 Other

9.3.1 Small Object Recognition

Small objects in isolation are easily detected by searching for clusters of ridge, river and edge contours, which is carried out with the texture analysis 1.3.4. The results of the texture analysis are loaded with the saliency file (Section 2.2.4). The first output argument, named `Txa` in our examples, contains a data structure called `Txa.B1b` containing the various bias statistics. In particular the bias for numerosity, [Num](#), and the one for high contrast outline often small objects in isolation. The corresponding candidates can be identified using the type variable `Typ`:

```
Bnum = Txa.B1b.Typ==1;      % numerous
Benk  = Txa.B1b.Typ==8;      % high-contrast
```

which can be used to access the bounding boxes in `Txa.B1b.Box`, ie.

```
BboxNum = Txa.Blb.Box( Bnum, : );
```

The example script `exsbPlotSalc.m` had already demonstrated this. Directory `/Demos` holds an example script `exsbSmlObjDet.m`, that demonstrates this for two displays made of small targets. The example makes use of the parameter adjusting the window size for texture integration (`--txws`).

The less isolated an object occurs in a scene, such as a drone flying near a tree silhouette, the more we need to know about its characteristics to discriminate it from its surround. We then build a classifier that operates on contour information, as a first step, that will help to eliminate unlikely candidates. We increase the specificity of the classifier by including also information from radial descriptors and perhaps even shape descriptors (`shp`), from focii, from shape extraction (`shpx`), or applying the entire descriptor extraction process to solely that image part.

If an object is heavily occluded - quasi camouflaged -, such as a traffic sign covered by leaves, then we need a Deep (Network) Detector, that excels at integrating, dispersed local information. Applying a Deep Detector to the entire image is of course expensive, and also struggles with low-contrast situations, but the methodology introduced so far enables to apply the detector more specifically and therefore to accelerate its use.

9.3.2 Anomaly and Change Detection

Anomaly Detection The segmentation process is predestined for anomaly detection, because it segments anything, irrespective of contrast and shape. A transparent, plastic bag floating across the road pavement is easily detected, as well as different pavement colors as pointed out already above (Road Surface). The curve partitions (arcs and straighters) also allow to understand the structure of the anomaly. If the location of the anomaly can be determined, it makes sense to apply the segregation process `sgrRGB` (Section 8) or to apply shape extraction (`shpx`).

Change Detection The entire feature extraction output is predestined for change detection due to its thorough topological analysis. The matching program `mvec1` can be deployed for finding potential changes (Section 3.1), ie. applied to focii; or the shape extraction and matching process (`shpx` and `mshp`).

9.3.3 Collecting Annotations

The output of the region segmentation process (Section 1.3.2) can be deployed to collect objects with clear boundaries. To understand that, it is best to run the demo programs in directory `/DemoBaum`. There exists one program for gray-level segmentation, called `baumgrau`, and another for chromatic segmentation `baumfarb`. The two program binaries can be run through example scripts `exsbBaum1grau.m` and `exsbBaum1rgb.m`, respectively.

Since many objects correspond to one or few regions, it is much easier to collect annotated material by clicking on regions, instead of outlining them with multiple clicks. The explicit specification of a complete bounding box is only necessary for colorful objects that appear in front of complex backgrounds, a situation that is equivalent to camouflage. Even in such situations, we gain sharp part and object boundaries. Using the program binary `sgrRGB`, one can obtain the sharpest region silhouettes, because it focuses on one RGB triplet by specification.

9.3.4 Active Vision

In active vision we are confronted with a larger range of visual signals than is present in photographs, because usually we photograph what is of interest and move the camera such that the object lies in the center of the image, and adjust the zoom appropriately. To efficiently mimick this phase of visual orienting, the saliency file can be deployed as starting information (Section 2.2.4). Its output can be used to gauge whether we have the appropriate zoom and if there is an object or scene part in the image center. We recall that we can load the `slc` file with:

```
[Txa Shp Ens Dsc] = LoadDescSalc( pthSalc );
```

Perhaps the first step is to distinguish between texture and non-texture, which we can be done with data structures `Txa` and `Dsc`, in particular `Txa.Gst` and `Dsc.MxRngRR`. We observe the ridge-river contrast value for the first level

```
Dsc.MxRngRR(1)      (or Dsc.MxRngRR[0] for zero-indexing as in C or Python)
```

in order to adjust the zoom. Simultaneously we observe the numerosity and blankness bias:

```
TxaFld.Gst.PrpPres(1) % numerosity
TxaFld.Gst.PrpPres(2) % blankness
```

The demo script `exsbSalBlobs.m` displays those values for two different images, one full of texture, the other void of it.

We would also observe the associated bounding boxes, see the code lines above for small object recognition (Section 9.3.1). In this case, we observe whether a blob of texture lies in the image center, e.g. foliage of a tree or plant, or an object containing a texture; conversely, if a blankness blob is present, it is likely to be an object or scene part of homogenous color. Either one might be better outlined by a shape, for which we therefore also observe the statistics in `Shp`. The structure `Ens` holds information that is better suited for an analysis of heterogeneous scenes.

9.4 Methodological Fusion

In certain scenarios, other methodologies outperform the present methodology - with its current parameter values. It therefore makes sense to combine them, as was suggested already previously. Since structure is well expressed and identified with the present methodology, the other methods can be deployed in a much more specific manner. Here we summarize potential approaches.

9.4.1 Local Features

Local features, such as histograms of gradient values, excel at identifying scenes in very large environments, ie. place recognition in the world. Those features are often sampled randomly from an image. With the contour and region features provided here, one can test a more directed approach to apply them, ie. taken at proposals as provided by the saliency file (`slc`).

9.4.2 Deep Networks/Learning

We firstly discuss the use of convolutional neural networks (CNN), that typically take pixels as input, also termed end-to-end learned. Then we mention the use of those networks that take processed input. Finally, there are possibilities to deploy Graph CNNs.

Pixels as Input (CNNs) CNNs are particularly good in two situations: for discriminating complex (multi-region) objects, or subtly different categories; and for the detection of heavily occluded objects, such as a traffic sign covered by tree branches, a situation also termed "camouflaged". Since CNNs require much resources, it makes sense to provide candidates with our structural approach in order to accelerate the recognition process, as discussed already with the pipelines (Section 1.6). This means that one would train the CNNs according to the confusions of our structural approach. For the case of object detection, one would provide candidate locations, where the object detector is applied specifically, in order to reduce its deployment.

Processed Input For networks that take vectors as input, such as transformers, there exists of course many possibilities to test our multi-dimensional attribute spaces. For networks that take tabular information or words as input, there exists equally many possibilities to generate them.

Graph CNNs Graph CNNs are useful for discriminating distributions in low-dimensional space. They could be applied to boundaries in order to discriminate subtly different shape silhouettes.

Appendix A

Feature Files

These files contain generic information about the features. Some of them were originally developed to make a comparison with object proposal studies. The `.Bbox` file contains the bounding boxes of regions in original map resolution. The `.BonBbox` contains the bounding boxes upscaled to original image resolution. The `.BonAsp` file contains some boundary aspects. The `.BonPix` file contains the boundary pixels.

A.1 Bounding Boxes (`.Bbox`)

The bounding boxes for regions can be saved by setting the (long option) flag `--saveBbox`. They are written in text format to a file with extension `.Bbox`. The list contains all bounding boxes. This allows to subselect according to the desired tasks.

The information in this file is minimal. More information is provided in the `.BonXXX` files (upcoming in Section A.3). But here we introduce the organization of the bounding boxes as well as the border values.

The first two integer values of the file hold the number of levels `nLev` and segmentation depth `depth` applied in the run. The following numbers hold the region count for each segmentation map, `nBbox`, saved looping levels as the outer loop and looping depth as the inner loop. The example below shows that for two levels and depth equal 3 (using zero-indexing).

```
nLev
depth
nBbox_Lev0_Depth0
nBbox_Lev0_Depth1
nBbox_Lev0_Depth2
nBbox_Lev1_Depth0
nBbox_Lev1_Depth1
nBbox_Lev1_Depth2
```

Then the bounding boxes follow. They are organized analogously to the above inner/outer loop: first all bounding boxes of `[lev=0,depth=0]`, then those of `[lev=0,depth=1]`, etc. A bounding box contains 6 parameters.

```
top, bottom, left, right, area, border      (for lev=0, depth=0)
top, bottom, left, right, area, border
...
top, bottom, left, right, area, border      (for lev=1, depth=2)
```

The parameters describe:

- **top, bottom, left, right**: absolute coordinates that correspond to the map size of the pyramidal level. Thus ones needs to upsample them by multiplying with the corresponding factor (2, 4, 8, ...).
- **area**: size of bounding box, calculated with the first four parameters.
- **border**: number of touches with the four image sides. The values are:

```

0          no touches: off border
1-4        at one border, directions NESW (top, rite, bot, left)
11-14      at two borders, directions NE,ES,SW,NW (topright, ...)
15,16     "   "   "   "   , NS axis, WE axis
101-3      at three borders
200        touching all borders

```

The bounding box sizes are typically slightly too small in comparison to annotations in datasets, partly due to the segmentation procedure and partly due to downsampling. Adding margins achieves better annotation correspondence, ie. margin values that correspond to the pyramid level.

Of course one can perform better selections with more information such as boundary contrast and perhaps region attributes, which is included in the **.BonXXX** files, upcoming in Section A.3.

A.2 Contour Endpoints (**.CntEpt**)

The contour endpoints are written to file if the long option **--saveBbox** is set - it is the same option as for the file generating the bounding boxes introduced above.

The points for the contour segments consist of the two endpoints as well as their midpoint. The points are written per level, per contour type and per point type. They are saved in binary format with extension **CntEpt**.

The first value holds the number of levels. Then each level of the pyramid is written separately with firstly the points of the ridge contours, then those of the river contours and eventually those of the edge contours. The points are written blockwise (and not rowwise as in case of the bounding boxes). The first value holds the number of descriptors. Then follow first all coordinates of the first endpoint (for that level); then all coordinates of the second endpoint; followed by all coordinates for the midpoint. The coordinates **coords** are saved as row/column pairs, per point.

```

nLev
nRdg  (# of ridge contours for lev=0)
[ridge coords of 1st endpoint for lev=0]
[ridge coords of 2nd endpoint for lev=0]
[ridge coords of midpoint      for lev=0]
nRiv  (# of river contours for lev=0)
[river coords of 1st endpoint for lev=0]
[river coords of 2nd endpoint for lev=0]
[river coords of midpoint      for lev=0]
nEdg  (# of edge contours for lev=0)
[edge  coords of 1st endpoint for lev=0]
[edge  coords of 2nd endpoint for lev=0]
[edge  coords of midpoint      for lev=0]
nRdg  (# of ridge contours for lev=1)
[ridge coords of 1st endpoint for lev=1]
[ridge coords of 2nd endpoint for lev=1]
[ridge coords of midpoint      for lev=1]
...

```

As with bounding boxes, the segment coordinates are absolute values corresponding to the map size of the pyramidal level. They need to be upsampled to match the

original image resolution if they are used as object/part proposals, and given some spatial width by adding some corresponding value.

A.3 Boundary Information (`.BonBbox`, `.BonAsp`, `.BonPix`)

If the flag `--saveBon` is set, then the full boundary information is saved to three separate files.

- `BonBbox`: bounding box information, similar to `Bbox` (as introduced above), but in different format.
- `BonAsp`: more boundary aspects.
- `BonPix`: boundary pixels.

The two scripts `exsbPlotBon.m` and `exsbPlotBonPix.m` demonstrate how to load those data files. The file with extension `BonBbox` contains essentially the same information as the `Bbox` file (described above under Section A.1), but in slightly different format and with additional information. The differences are:

- the bounding boxes are concatenated across depth, but the depth information is still available as the 6th parameter.
- the box coordinates are scaled to original size already.
- the contrast value for the boundaries is given as 5th parameter.

The file with extension `BonAsp` lists some additional boundary aspects:

- chromatic values red, green, blue for the pixels along the boundary, not the region inside. Thus for small regions it may not be an optimal chromatic representation.
- coverage of the boundary area, as proportion of the image/map
- border values as introduced under Section A.1 already.
- perimeter, which is given as absolute value.
- area in pixels for the connected component (not the boundary itself), thus excluding holes
- area of boundary and thus including holes.

The `BonPix` file is loaded with routine `LoadBonPix.m`.

Appendix B

Notation

Our code notation is leaned toward the Java notation that uses concatenated, capitalized words and syllables. The underscore sign '_' is not used for variables but for function names, where a prefix denotes the type (or class) of a function or routine. For example a prefix made of a single letter, such as `f_`, `i_`, `u_`, stands for computation, initialization and utility, respectively. But also prefixes of multiple letters are used. Only few routines are named without any underscore, that then contain a verb (or syllable of a verb) to express 'action' (ie. `Load`, `Save`, `Read`, ...) in order to distinguish themselves from variables.

- `f_Func`: routines starting with `f_` compute important functionality, such as feature extraction, feature manipulation, etc. The function name is composed of 'syllables' of three to four letters, aligned from abstract to more detailed.
- `i_Func`: routines starting with `i_` initialize an algorithm, process, etc..
- `u_Func`: functions starting with `u_` are utility functions carrying out administration, support, etc.
- `p_Func`: functions starting with `p_` are plotting routines.
- `LoadX`: loads from file with path being specified as function argument. Such functions typically call `Read` routines, as explained below.
- `SaveX`: saves data to file with path being specified as function argument. Such functions typically call `Write` routines, as explained below.
- `ReadX`: reads from file with filepointer given as function argument. They are usually called from a loading script `LoadX` (see above).
- `WriteX`: writes to file with filepointer given as function argument. They are usually called from a saving script `SaveX` (see above).
- `PlotX`: comprises a longer list of plotting instructions, calling usually plotting routines `p_Func`.
- `RennX`: runs a program binary/executable from Matlab using the Matlab function `dos`. It is a wrapper routine facilitating the use of the binary with its options, e.g. script `RennDscx.m` runs program `dscx`.
- `exsbX`: is a demonstration script showing how to apply a program, e.g. `exsbDscx.m` runs program `dscx`. These scripts usually call the corresponding wrapper routine `RennX` as an example.

The following notation is used for variables:

nX	number of X, e.g. nLev
stcX	structure of X
szX	size of X, e.g. szV , szH , vertical, horizontal size

For more explanations on the notation see also my Computer Vision overview:

<https://www.researchgate.net/publication/336460083>